

```
In [1]: # Initialize Otter
import otter
grader = otter.Notebook("lab4.ipynb")

In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.linear_model import Ridge, LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.tree import DecisionTreeRegressor
from sklearn.svm import SVR
from src.utils import custom_train_test_split, compute_feature_engineering, get_im
from sklearn.pipeline import make_pipeline
from sklearn.compose import make_column_transformer
from sklearn.preprocessing import OneHotEncoder, StandardScaler, PolynomialFeatures
from sklearn.metrics import mean_squared_error, r2_score, root_mean_squared_error
import shap
import os

DATA_PATH = "data/AB_NYC_2019.csv"
TARGET = "reviews_per_month"
RANDOM_STATE = 42
np.random.seed(RANDOM_STATE)
```

## Lab 4: Putting it all together in a mini project

For this lab, **you can choose to work alone or in a group of up to three students**. You are in charge of how you want to work and who you want to work with. Maybe you really want to go through all the steps of the ML process yourself or maybe you want to practice your collaboration skills, it is up to you! Just remember to indicate who your group members are (if any) when you submit on Gradescope. If you choose to work in a group, you only need to use one GitHub repo (you can create one on [github.ubc.ca](https://github.com) and set the visibility to "public"). If it takes a prohibitively long time to run any of the steps on your laptop, it is OK if you sample the data to reduce the runtime, just make sure you write a note about this.

### Submission instructions

`rubric={mechanics}`

You receive marks for submitting your lab correctly, please follow these instructions:

- Follow the general lab instructions.
- [Click here](#) to view a description of the rubrics used to grade the questions

- Make at least three commits.
- Push your `.ipynb` file to your GitHub repository for this lab and upload it to Gradescope.
  - Before submitting, make sure you restart the kernel and rerun all cells.
- Make sure to only make one gradescope submission per group, and to assign all group members on gradescope at submission time.
- Also upload a `.pdf` export of the notebook to facilitate grading of manual questions (preferably WebPDF, you can select two files when uploading to gradescope)
- Don't change any variable names that are given to you, don't move cells around, and don't include any code to install packages in the notebook.
- The data you download for this lab **SHOULD NOT BE PUSHED TO YOUR REPOSITORY** (there is also a `.gitignore` in the repo to prevent this).
- Include a clickable link to your GitHub repo for the lab just below this cell
  - It should look something like this [https://github.ubc.ca/MDS-2020-21/DSCI\\_531\\_labX\\_yourcwl](https://github.ubc.ca/MDS-2020-21/DSCI_531_labX_yourcwl).

*Points:* 2

FIND HERE OUR REPOSITORY 

**Collaborators:**

- Elí González Zequeida (`@Eligozo75`): Section 1
- Héctor Palafox Prieto (`@hpalafoxp`): Section 1
- Luis Alonso Álvarez Portugal (`@luisalonso8`): Section 2

## Introduction

In this lab you will be working on an open-ended mini-project, where you will put all the different things you have learned so far in 571 and 573 together to solve an interesting problem.

A few notes and tips when you work on this mini-project:

### Tips

1. Since this mini-project is open-ended there might be some situations where you'll have to use your own judgment and make your own decisions (as you would be doing when you work as a data scientist). Make sure you explain your decisions whenever necessary.
2. **Do not include everything you ever tried in your submission** -- it's fine just to have your final code. That said, your code should be reproducible and well-documented. For example, if you chose your hyperparameters based on some hyperparameter optimization experiment, you should leave in the code for that experiment so that

someone else could re-run it and obtain the same hyperparameters, rather than mysteriously just setting the hyperparameters to some (carefully chosen) values in your code.

3. If you realize that you are repeating a lot of code try to organize it in functions. Clear presentation of your code, experiments, and results is the key to be successful in this lab. You may use code from lecture notes or previous lab solutions with appropriate attributions.

## Assessment

We don't have some secret target score that you need to achieve to get a good grade. **You'll be assessed on demonstration of mastery of course topics, clear presentation, and the quality of your analysis and results.** For example, if you just have a bunch of code and no text or figures, that's not good. If you instead try several reasonable approaches and you have clearly motivated your choices, but still get lower model performance than your friend, don't sweat it.

## A final note

Finally, the style of this "project" is different from other assignments. It'll be up to you to decide when you're "done" -- in fact, this is one of the hardest parts of real projects. But please don't spend WAY too much time on this... perhaps "several hours" but not "many hours" is a good guideline for a high quality submission. Of course if you're having fun you're welcome to spend as much time as you want! But, if so, try not to do it out of perfectionism or getting the best possible grade. Do it because you're learning and enjoying it. Students from the past cohorts have found such kind of labs useful and fun and we hope you enjoy it as well.

## 1. Pick your problem and explain the prediction problem

`rubric={reasoning}`

In this mini project, you will pick one of the following problems:

1. A classification problem of predicting whether a credit card client will default or not. For this problem, you will use Default of Credit Card Clients Dataset. In this data set, there are 30,000 examples and 24 features, and the goal is to estimate whether a person will default (fail to pay) their credit card bills; this column is labeled "default.payment.next.month" in the data. The rest of the columns can be used as features. You may take some ideas and compare your results with the associated research paper, which is available through the UBC library.

OR

2. A regression problem of predicting `reviews_per_month`, as a proxy for the popularity of the listing with New York City Airbnb listings from 2019 dataset. Airbnb could use this sort of model to predict how popular future listings might be before they are posted, perhaps to help guide hosts create more appealing listings. In reality they might instead use something like vacancy rate or average rating as their target, but we do not have that available here.

**Your tasks:**

1. Spend some time understanding the problem and what each feature means. Write a few sentences on your initial thoughts on the problem and the dataset.
2. Download the dataset and read it as a pandas dataframe.
3. Carry out any preliminary preprocessing, if needed (e.g., changing feature names, handling of NaN values etc.)

*Points:* 3

Initially, problem 2 is an interesting topic, but perhaps not the best features or target to predict. Reviews per month can serve as an acceptable proxy for predicting the popularity of future listings, though in reality, Airbnb might prefer metrics like vacancy rate or average rating as targets if they were available. Either way, `reviews_per_month` can serve as a reasonable target for this regression problem. Some features seem promising based on preliminary insights:

1. Neighborhood is an important variable, as people often consider location when choosing where to stay and may be more likely to write reviews for listings in desirable areas with good subway or public transportation access.
2. Price could provide the most predictive power to the model (as it is in many times, the most important variable when a person choose a place at Airbnb and write a review if it was a competitive price and a good place).
3. Availability could also be useful, as there could be "good places in photos" or with good reviews but if there is enough availability, people can not write a review or choose it as good place to stay.
4. However, number of reviews is probably highly correlated with the target variable, which could potentially lead to data leakage.
5. Room type may also affect the visitor experience and influence review behavior, but maybe in a smaller magnitude than other features.
6. Calculated host listing counts: here we can see a difference between people that have professional or experience with Airbnb compared to new ones that are just beginning to list a property.
7. Geographic features may be also a factor as people intend to be nearer tourist attractions as Central Park, Wall Street, Manhattan or Times Square than in other places.

8. Booking constraints can deter some guests or attracts longer travelers (who maybe have higher income).

Additionally, there are challenges with the review data itself: many people stay in a place and enjoy it but don't write a review, while dissatisfied guests are far more likely to leave feedback, which can create imbalance in the data and also a clear selection bias. Finally, missing values are likely present and will need to be addressed and handled correctly. All of this is preliminary, even before looking into the data, splitting into the train/test and performing an EDA.

In [ ]:

## 2. Data splitting

rubric={reasoning}

### Your tasks:

1. Split the data into train and test portions.

Make the decision on the `test_size` based on the capacity of your laptop.

Points: 1

In [3]:

```
path = 'data/split_data/train_set.csv'

# Check if the path is a file
if os.path.isfile(path):
    print(f"The file exists.")
else:
    print("Didn't find the datasets. Creating them from scratch:")
    %run src/jobs/01_split_data.py
```

The file exists.

In [4]:

```
df = pd.read_csv(DATA_PATH)
df = df.dropna(subset=[TARGET])
print(df.shape)
df.head()
```

(38843, 16)

Out[4]:

	<b>id</b>	<b>name</b>	<b>host_id</b>	<b>host_name</b>	<b>neighbourhood_group</b>	<b>neighbourhood</b>	<b>latitude</b>
<b>0</b>	2539	Clean & quiet apt home by the park	2787	John	Brooklyn	Kensington	40.64749
<b>1</b>	2595	Skylit Midtown Castle	2845	Jennifer	Manhattan	Midtown	40.75362
<b>3</b>	3831	Cozy Entire Floor of Brownstone	4869	LisaRoxanne	Brooklyn	Clinton Hill	40.68514
<b>4</b>	5022	Entire Apt: Spacious Studio/Loft by central park	7192	Laura	Manhattan	East Harlem	40.79851
<b>5</b>	5099	Large Cozy 1 BR Apartment In Midtown East	7322	Chris	Manhattan	Murray Hill	40.74767

In [5]:

```
df_train, df_test = custom_train_test_split(df, "host_id")

print("PCT of rows in train:", 100.0 * len(df_train) / len(df))
print("PCT of rows in test:", 100.0 * len(df_test) / len(df))
```

PCT of rows in train: 69.99974255335582  
PCT of rows in test: 30.000257446644184

In [6]:

```
not_use_cols = [
    'id',
    'host_id',
    'host_name',
    'last_review',
]
use_cols = [c for c in df if c not in not_use_cols]
df_train = df_train[use_cols].copy()
df_test = df_test[use_cols].copy()
```

In [7]:

```
df_train.to_csv("data/split_data/train_set.csv", index=False)
df_test.to_csv("data/split_data/test_set.csv", index=False)
```

### 3. EDA

rubric={viz,reasoning}

Perform exploratory data analysis on the train set.

**Your tasks:**

1. Include at least two summary statistics and two visualizations that you find useful, and accompany each one with a sentence explaining it.
2. Summarize your initial observations about the data.
3. Pick appropriate metric/metrics for assessment.

*Points: 6*

First, it is important to note that we split the data into train and test sets (70-30) by host\_id to prevent data leakage. This ensures that all listings from the same host are in either the train or test set, but not both. This way, the model learns to generalize to new hosts rather than just new listings from already known hosts.

EDA

a) After splitting, both categorical and numeric data types are clean with no missing values in the training set (27,190 entries). We removed 20% of the original data that had missing reviews\_per\_month values (corresponding to listings with zero reviews). The target variable distribution in the training set shows: median is 0.72 reviews per month, the minimum is 0.01 "reviews per month", the mean is 1.37 reviews and the maximum is 58.5 reviews per month. So, we are going to model "reviews per month" with reviews that have, at least, one review so it effectively predict review frequency for listings that attract visitors, rather than if a listing receive a review, at all.

b) Geographic Patterns in reviews: The mean of "reviews\_per\_month" is much higher in outer boroughs than in NYC's tourist centers. Counterintuitively, outer boroughs (Queens: 1.96, Bronx: 1.76) have higher average reviews per month than tourist centers (Manhattan: 1.26, Brooklyn: 1.28). Possible explanations include: (a) budget travelers stay longer and are more likely to leave reviews, (b) sample size bias in outer boroughs where only the best listings survive, or (c) tourists in Manhattan prefer hotels over Airbnb, as they may be higher-income travelers.

c) Price Outliers: We observe very large outliers across many numeric features, with right-skewed distributions. Price ranges up to 10000 USD.

d) Additional Outliers Across Features: Several other features contain extreme outliers:

- One listing requires a minimum stay of 999 days
- Maximum "reviews\_per\_month" is 58.5, equivalent to nearly two reviews per day, which is highly unlikely
- Maximum total reviews is 629, which seems unreasonably high
- One host has up to 232 listings

- Some listings show 365 days of availability, which is unrealistic for popular properties
- e) Neighborhood-Level Patterns: At the neighborhood level, there are no clear trends. Additionally, some neighborhoods with high mean reviews have very few listings, suggesting these patterns may not be reliable or generalizable.
- f) Room Type Analysis: Private rooms and shared rooms (1.44 and 1.43 reviews/month, respectively) have slightly higher average reviews\_per\_month than entire homes/apartments (1.30). However, the dataset contains more listings in Brooklyn and Manhattan (44% and 41%, respectively) than in the other boroughs. We hypothesize that private rooms attract budget travelers who are more likely to leave reviews compared to higher-income travelers who book entire homes.
- g) Geographic Distributions: Latitude and longitude show approximately normal distributions within reasonable ranges for NYC.
- h) Spatial Relationships: Visualizations reveal mostly no clear spatial trends between geographic features and reviews\_per\_month. The only strong relationship observed is between "number\_of\_reviews" and "reviews\_per\_month", which is expected but indicates potential data leakage if "number\_of\_reviews" is used as a feature.

```
In [8]: df_train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 27190 entries, 0 to 27189
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
 ---  --  
 0   name             27184 non-null    object  
 1   neighbourhood_group  27190 non-null    object  
 2   neighbourhood      27190 non-null    object  
 3   latitude          27190 non-null    float64 
 4   longitude         27190 non-null    float64 
 5   room_type         27190 non-null    object  
 6   price             27190 non-null    int64  
 7   minimum_nights    27190 non-null    int64  
 8   number_of_reviews  27190 non-null    int64  
 9   reviews_per_month  27190 non-null    float64 
 10  calculated_host_listings_count 27190 non-null    int64  
 11  availability_365   27190 non-null    int64  
dtypes: float64(3), int64(5), object(4)
memory usage: 2.5+ MB
```

```
In [9]: df_train.isna().mean()
```

```
Out[9]: name          0.000221  
neighbourhood_group 0.000000  
neighbourhood       0.000000  
latitude            0.000000  
longitude           0.000000  
room_type           0.000000  
price               0.000000  
minimum_nights      0.000000  
number_of_reviews   0.000000  
reviews_per_month   0.000000  
calculated_host_listings_count 0.000000  
availability_365    0.000000  
dtype: float64
```

```
In [10]: df_train.describe()
```

	latitude	longitude	price	minimum_nights	number_of_reviews	rev
<b>count</b>	27190.000000	27190.000000	27190.000000	27190.000000	27190.000000	
<b>mean</b>	40.727681	-73.950990	142.489702	5.790953	29.232953	
<b>std</b>	0.055250	0.046678	206.368954	16.108038	48.545982	
<b>min</b>	40.508680	-74.244420	0.000000	1.000000	1.000000	
<b>25%</b>	40.688240	-73.982270	68.000000	1.000000	3.000000	
<b>50%</b>	40.721180	-73.954550	100.000000	2.000000	9.000000	
<b>75%</b>	40.762910	-73.935342	169.000000	4.000000	33.000000	
<b>max</b>	40.908040	-73.712990	10000.000000	999.000000	629.000000	

```
In [11]: df_train.groupby("neighbourhood_group")[TARGET].mean().sort_values(ascending=False)
```

```
Out[11]: neighbourhood_group  
Queens          1.958524  
Staten Island   1.854872  
Bronx           1.761455  
Brooklyn        1.283295  
Manhattan       1.260512  
Name: reviews_per_month, dtype: float64
```

```
In [12]: df_train["neighbourhood"].value_counts(True)
```

```
Out[12]: neighbourhood
Williamsburg      0.083854
Bedford-Stuyvesant 0.083339
Harlem            0.056933
Bushwick           0.048547
Upper West Side    0.038911
...
West Farms        0.000037
Westerleigh        0.000037
Shore Acres         0.000037
Lighthouse Hill    0.000037
Silver Lake         0.000037
Name: proportion, Length: 216, dtype: float64
```

```
In [13]: df_train.groupby("neighbourhood", as_index=False).agg(
    {TARGET: ["mean", "count"]}
).sort_values(by=("reviews_per_month", "mean"), ascending=False)
```

```
Out[13]: neighbourhood  reviews_per_month
```

		mean	count
139	New Dorp Beach	5.500000	2
173	Silver Lake	5.490000	1
58	East Elmhurst	5.150650	123
208	Whitestone	5.060000	2
179	Springfield Gardens	4.743731	67
...	...	...	...
96	Holliswood	0.310000	1
145	Oakwood	0.285000	2
172	Shore Acres	0.280000	1
204	West Farms	0.160000	1
119	Marble Hill	0.138571	7

216 rows × 3 columns

```
In [14]: df_train["room_type"].value_counts(True)
```

```
Out[14]: room_type
Entire home/apt    0.519934
Private room       0.457595
Shared room         0.022471
Name: proportion, dtype: float64
```

```
In [15]: df_train.groupby("room_type")[TARGET].mean().sort_values(ascending=False)
```

```
Out[15]: room_type
Private room      1.445947
Shared room       1.433863
Entire home/apt  1.297404
Name: reviews_per_month, dtype: float64
```

```
In [16]: df_train["neighbourhood_group"].value_counts(True).sort_values(ascending=True)
```

```
Out[16]: neighbourhood_group
Staten Island    0.008606
Bronx           0.021478
Queens          0.117837
Manhattan        0.421736
Brooklyn         0.430342
Name: proportion, dtype: float64
```

```
In [17]: cont_cols = [c for c in df_train.describe()]

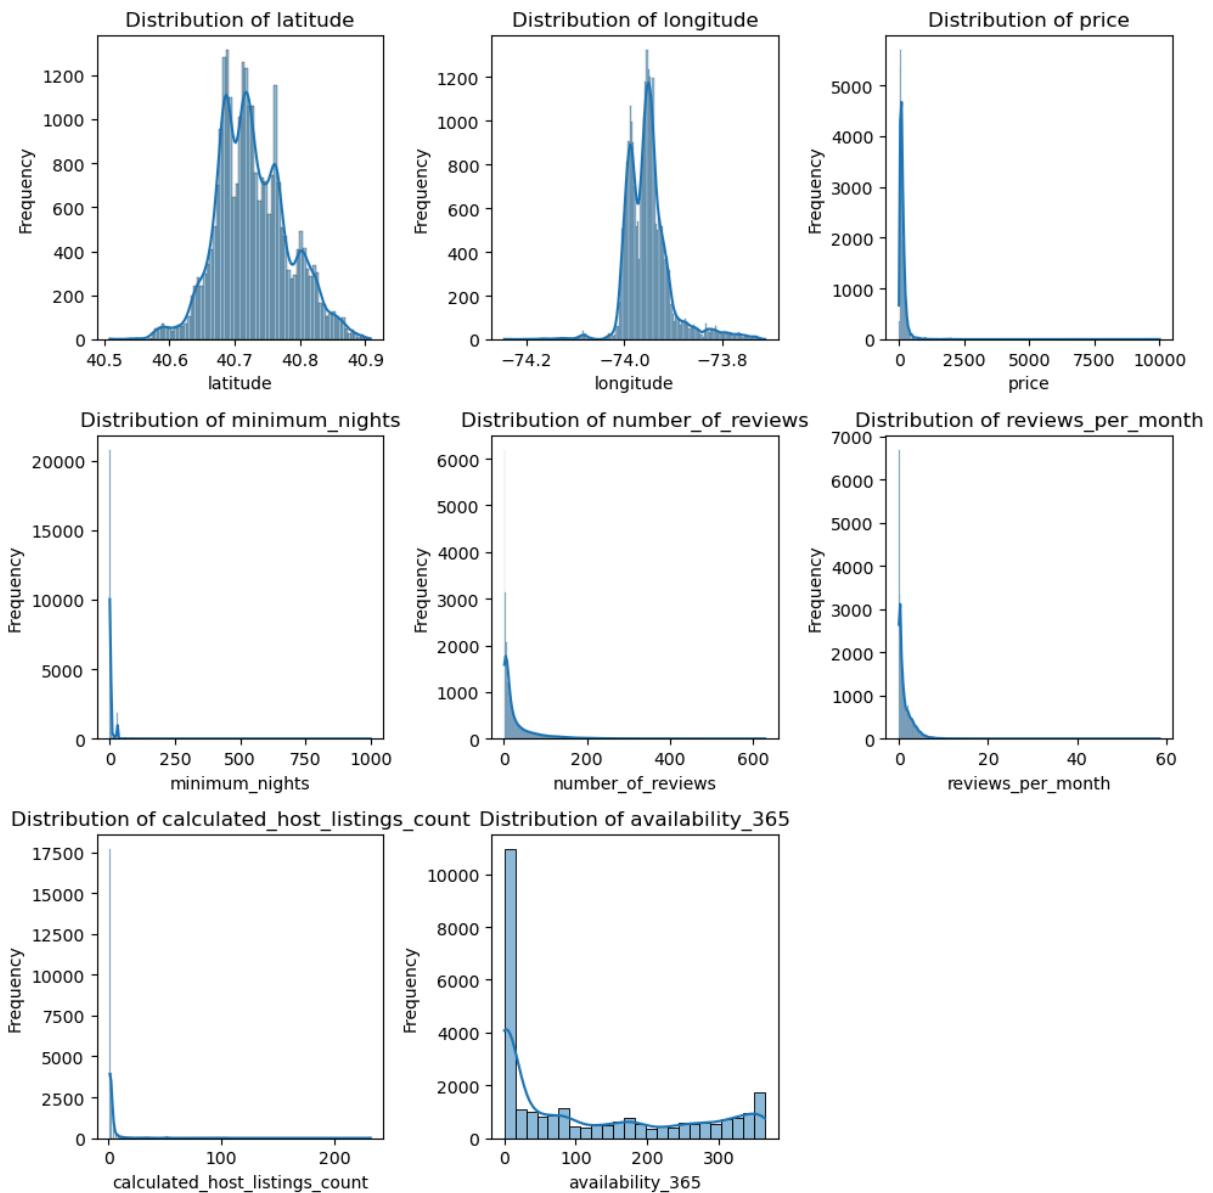
num_cols_per_row = 3
num_rows = int(np.ceil(len(cont_cols) / num_cols_per_row))

# Create the figure and subplots
fig, axes = plt.subplots(num_rows, num_cols_per_row, figsize=(10, 10))
axes = axes.flatten() # Flatten the 2D array of axes for easier iteration

# Plot histograms for each column
for i, col in enumerate(cont_cols):
    sns.histplot(df_train[col], kde=True, ax=axes[i])
    axes[i].set_title(f"Distribution of {col}")
    axes[i].set_xlabel(col)
    axes[i].set_ylabel("Frequency")

# Remove any unused subplots if the number of columns is not a multiple of 3
for j in range(len(cont_cols), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```



```
In [18]: num_cols_per_row = 3
num_rows = int(np.ceil(len(cont_cols) / num_cols_per_row))

# Create the figure and subplots
fig, axes = plt.subplots(num_rows, num_cols_per_row, figsize=(10, 10))
axes = axes.flatten() # Flatten the 2D array of axes for easier iteration

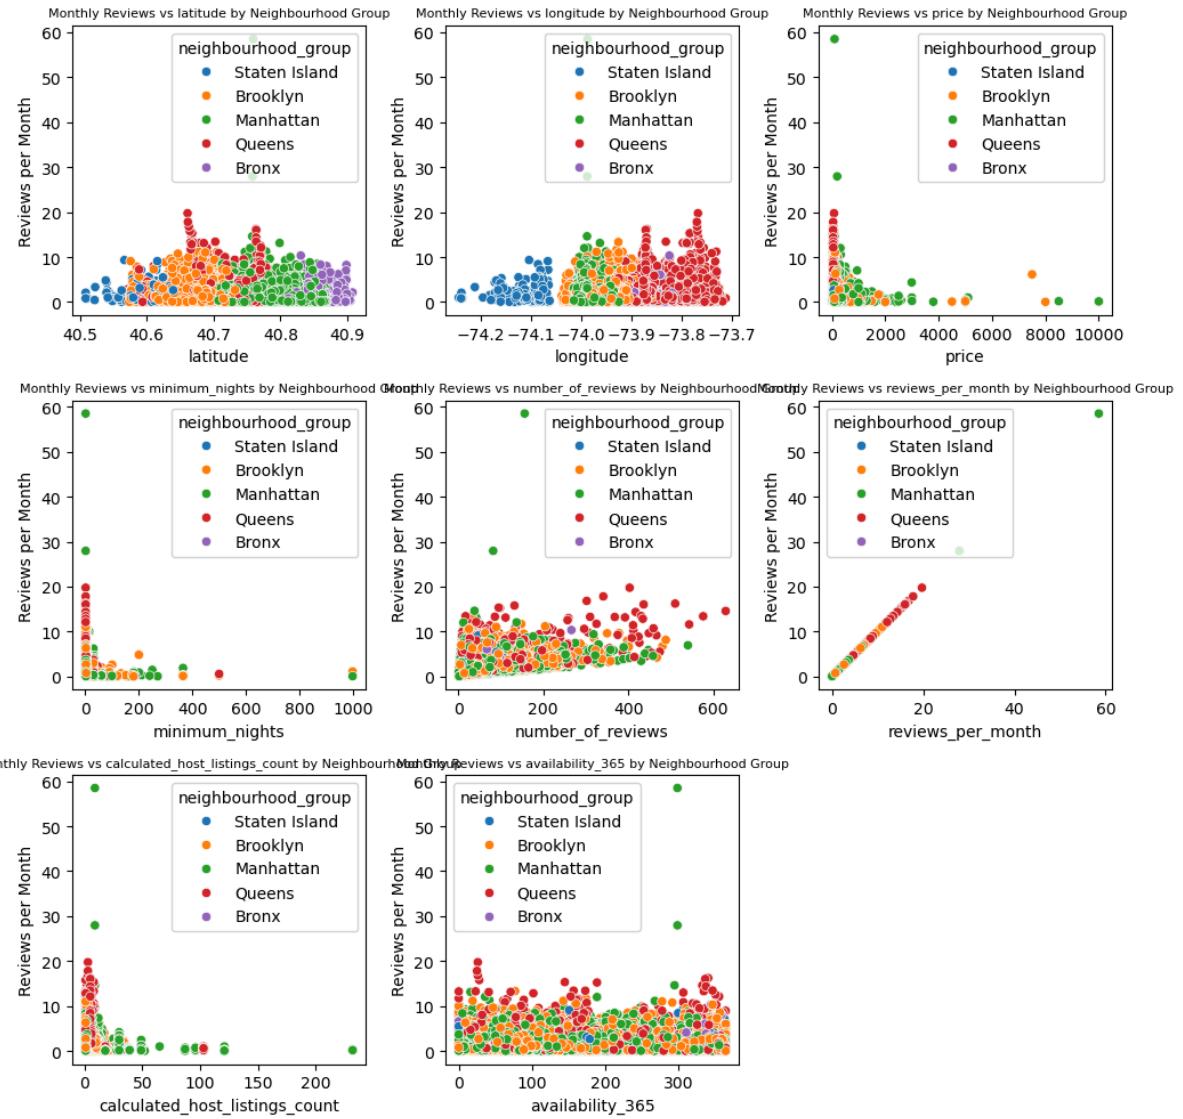
# Plot histograms for each column
for i, col in enumerate(cont_cols):
    sns.scatterplot(
        x=df_train[col],
        y=df_train[TARGET],
        hue=df_train["neighbourhood_group"],
        ax=axes[i],
    )
    axes[i].set_title(f"Monthly Reviews vs {col} by Neighbourhood Group", fontsize=12)
    axes[i].set_xlabel(col)
    axes[i].set_ylabel("Reviews per Month")
```

```

# Remove any unused subplots if the number of columns is not a multiple of 3
for j in range(len(cont_cols), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```



```

In [19]: num_cols_per_row = 3
num_rows = int(np.ceil(len(cont_cols) / num_cols_per_row))

# Create the figure and subplots
fig, axes = plt.subplots(num_rows, num_cols_per_row, figsize=(10, 10))
axes = axes.flatten() # Flatten the 2D array of axes for easier iteration

# Plot histograms for each column
for i, col in enumerate(cont_cols):
    sns.scatterplot(
        x=df_train[col],
        y=df_train[TARGET],
        hue=df_train["room_type"],
        ax=axes[i],
    )
)

```

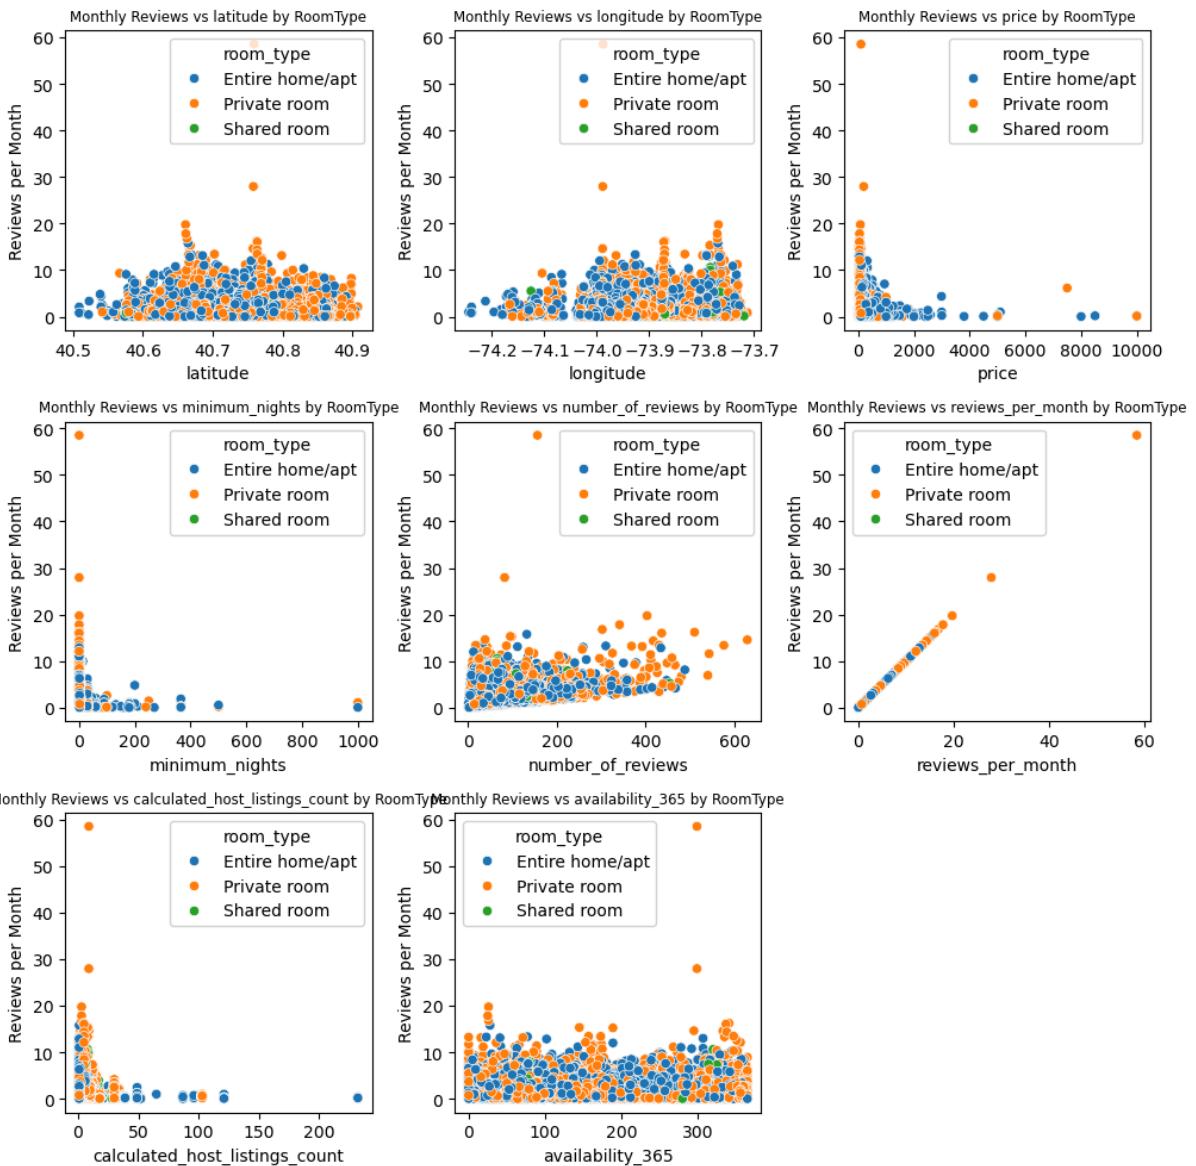
```

        axes[i].set_title(f"Monthly Reviews vs {col} by RoomType", fontsize=8.5)
        axes[i].set_xlabel(col)
        axes[i].set_ylabel("Reviews per Month")

# Remove any unused subplots if the number of columns is not a multiple of 3
for j in range(len(cont_cols), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```



```

In [20]: num_cols_per_row = 3
num_rows = int(np.ceil(len(cont_cols) / num_cols_per_row))

# Create the figure and subplots
fig, axes = plt.subplots(num_rows, num_cols_per_row, figsize=(10, 10))
axes = axes.flatten() # Flatten the 2D array of axes for easier iteration

# Plot histograms for each column
for i, col in enumerate(cont_cols):

```

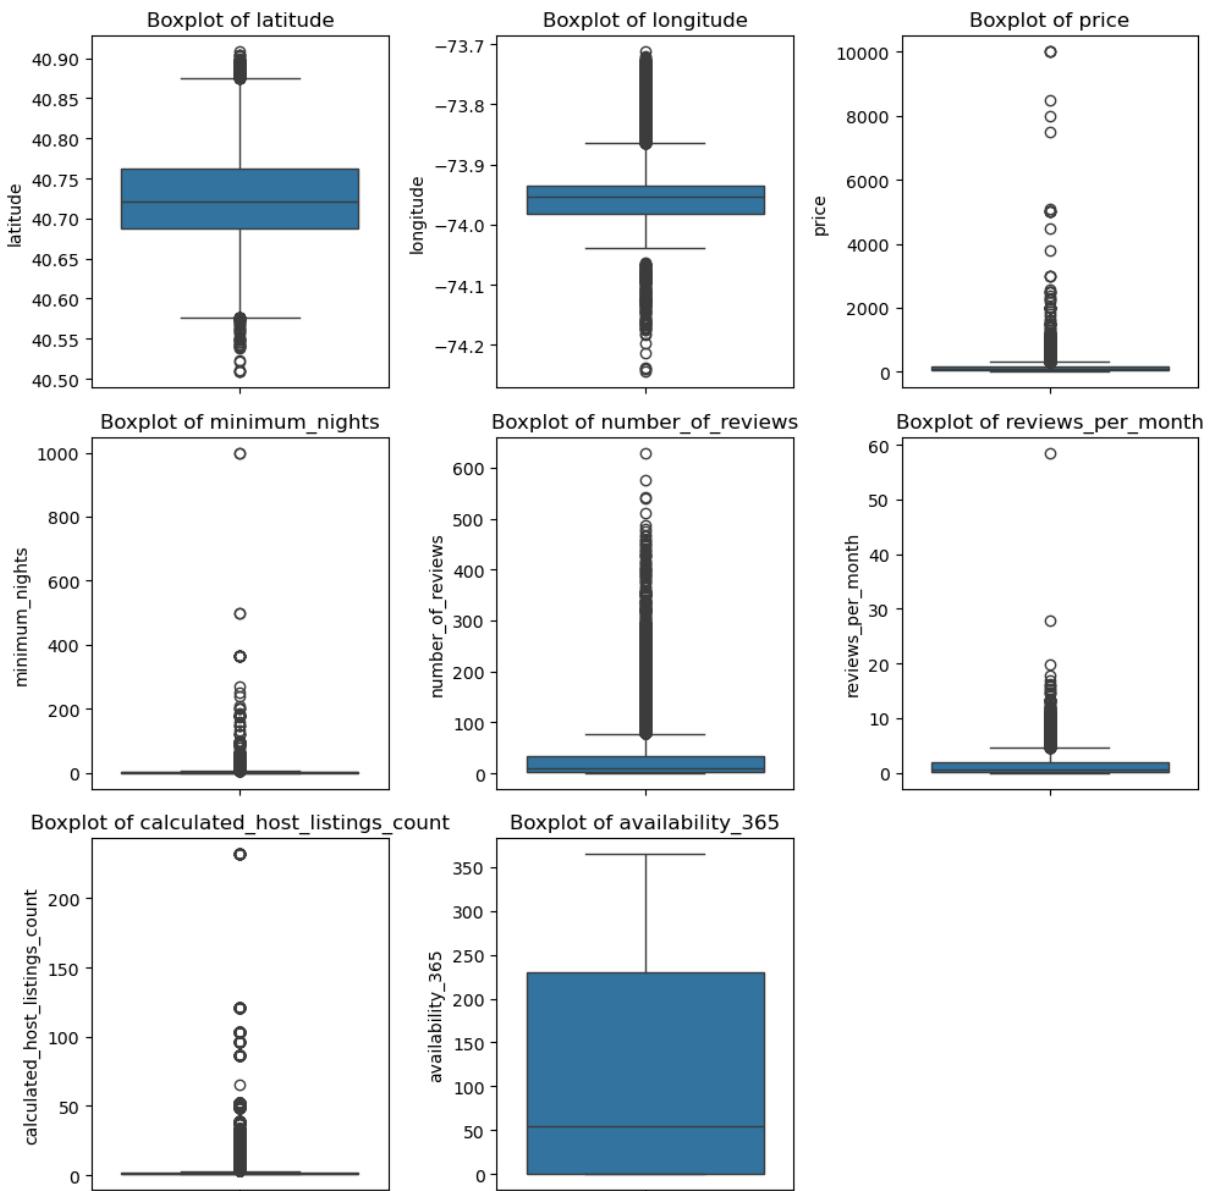
```

sns.boxplot(df_train[col], ax=axes[i])
axes[i].set_title(f"Boxplot of {col}")

# Remove any unused subplots if the number of columns is not a multiple of 3
for j in range(len(cont_cols), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()

```



```

In [21]: plt.rcParams.update({
    "font.size": 8,
    "axes.titlesize": 9,
    "axes.labelsize": 8,
    "xtick.labelsize": 7,
    "ytick.labelsize": 7,
})

num_cols = df_train.select_dtypes(include="number").columns.tolist()
cont_cols = [c for c in num_cols if c != TARGET]

```

```

method = "pearson"

corr_df = df_train[cont_cols + [TARGET]].corr(method=method)

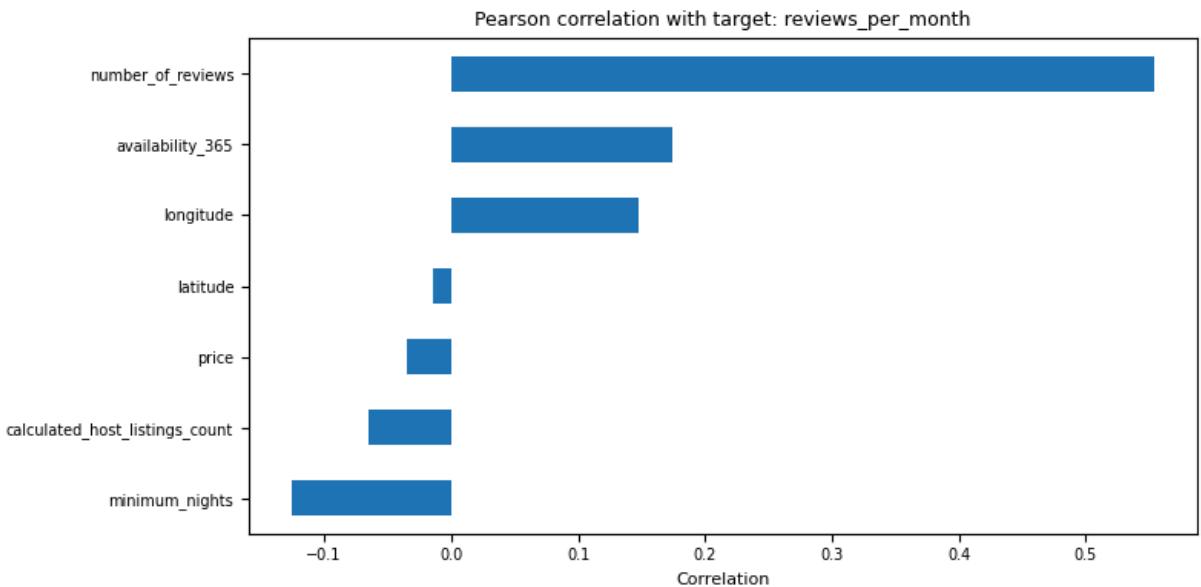
feat_corr = corr_df[TARGET].drop(TARGET).sort_values()

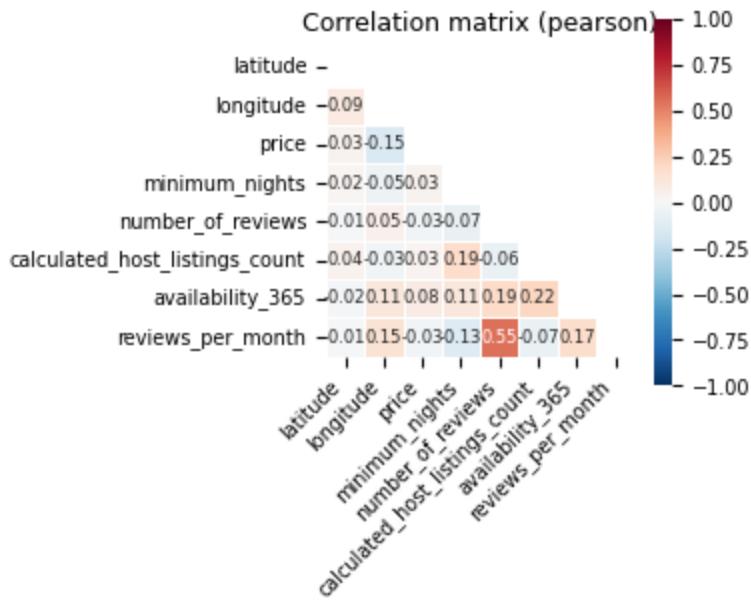
plt.figure(figsize=(8, max(4, 0.28 * len(feat_corr))))
feat_corr.plot(kind="barh")
plt.title(f"{{method.title()}} correlation with target: {{TARGET}}", fontsize=9)
plt.xlabel("Correlation", fontsize=8)
plt.yticks(fontsize=7)
plt.xticks(fontsize=7)
plt.tight_layout()
plt.show()

mask = np.triu(np.ones_like(corr_df, dtype=bool))

plt.figure(figsize=(1 + 0.35 * len(corr_df.columns), 1 + 0.35 * len(corr_df.columns))
sns.heatmap(
    corr_df,
    mask=mask,
    center=0,
    vmin=-1, vmax=1,
    cmap="RdBu_r",
    square=True,
    linewidths=0.5,
    annot=True,
    fmt=".2f",
    annot_kws={"size": 6},
    cbar_kws={"shrink": 0.7}
)
plt.title(f"Correlation matrix ({method})", fontsize=9)
plt.xticks(fontsize=7, rotation=45, ha="right")
plt.yticks(fontsize=7)
plt.tight_layout()
plt.show()

```





## 4. Feature engineering (Challenging)

`rubric={reasoning}`

**Your tasks:**

- Carry out feature engineering. In other words, extract new features relevant for the problem and work with your new feature set in the following exercises. You may have to go back and forth between feature engineering and preprocessing.

*Points: 0.5*

```
In [22]: df_train = pd.read_csv("data/split_data/train_set.csv")
df_test = pd.read_csv("data/split_data/test_set.csv")
```

```
In [23]: df_train = compute_feature_engineering(df_train)
print("\n\nFinished TRAIN feature engineering...\n")
df_test = compute_feature_engineering(df_test)
print("\n\nFinished TEST feature engineering...")
```

```
COMPUTING FEATURE: nyc_tm_distance_km  
COMPUTING FEATURE: nyc_cp_distance_km  
COMPUTING FEATURE: nyc_fd_distance_km  
COMPUTING FEATURE: jfk_airport_distance_km  
COMPUTING FEATURE: lga_airport_distance_km  
COMPUTING FEATURE: ewr_airport_distance_km  
COMPUTING FEATURE: avg_airport_distance_km  
COMPUTING FEATURE: avg_airport_distance_km
```

Finished TRAIN feature engineering...

```
COMPUTING FEATURE: nyc_tm_distance_km  
COMPUTING FEATURE: nyc_cp_distance_km  
COMPUTING FEATURE: nyc_fd_distance_km  
COMPUTING FEATURE: jfk_airport_distance_km  
COMPUTING FEATURE: lga_airport_distance_km  
COMPUTING FEATURE: ewr_airport_distance_km  
COMPUTING FEATURE: avg_airport_distance_km  
COMPUTING FEATURE: avg_airport_distance_km
```

Finished TEST feature engineering...

```
In [24]: new_features = [  
    "nyc_tm_distance_km",  
    "nyc_cp_distance_km",  
    "nyc_fd_distance_km",  
    "jfk_airport_distance_km",  
    "lga_airport_distance_km",  
    "ewr_airport_distance_km",  
    "avg_airport_distance_km",  
    "total_min_cost",  
    "len_name",  
    "nb_adj_in_name",  
    "nb_adv_in_name",  
    "nb_nouns_in_name",  
    "nb_propn_in_name",
```

```
"rate_adj_in_name",
"rate_adv_in_name",
"rate_nouns_in_name",
"rate_propn_in_name",
]

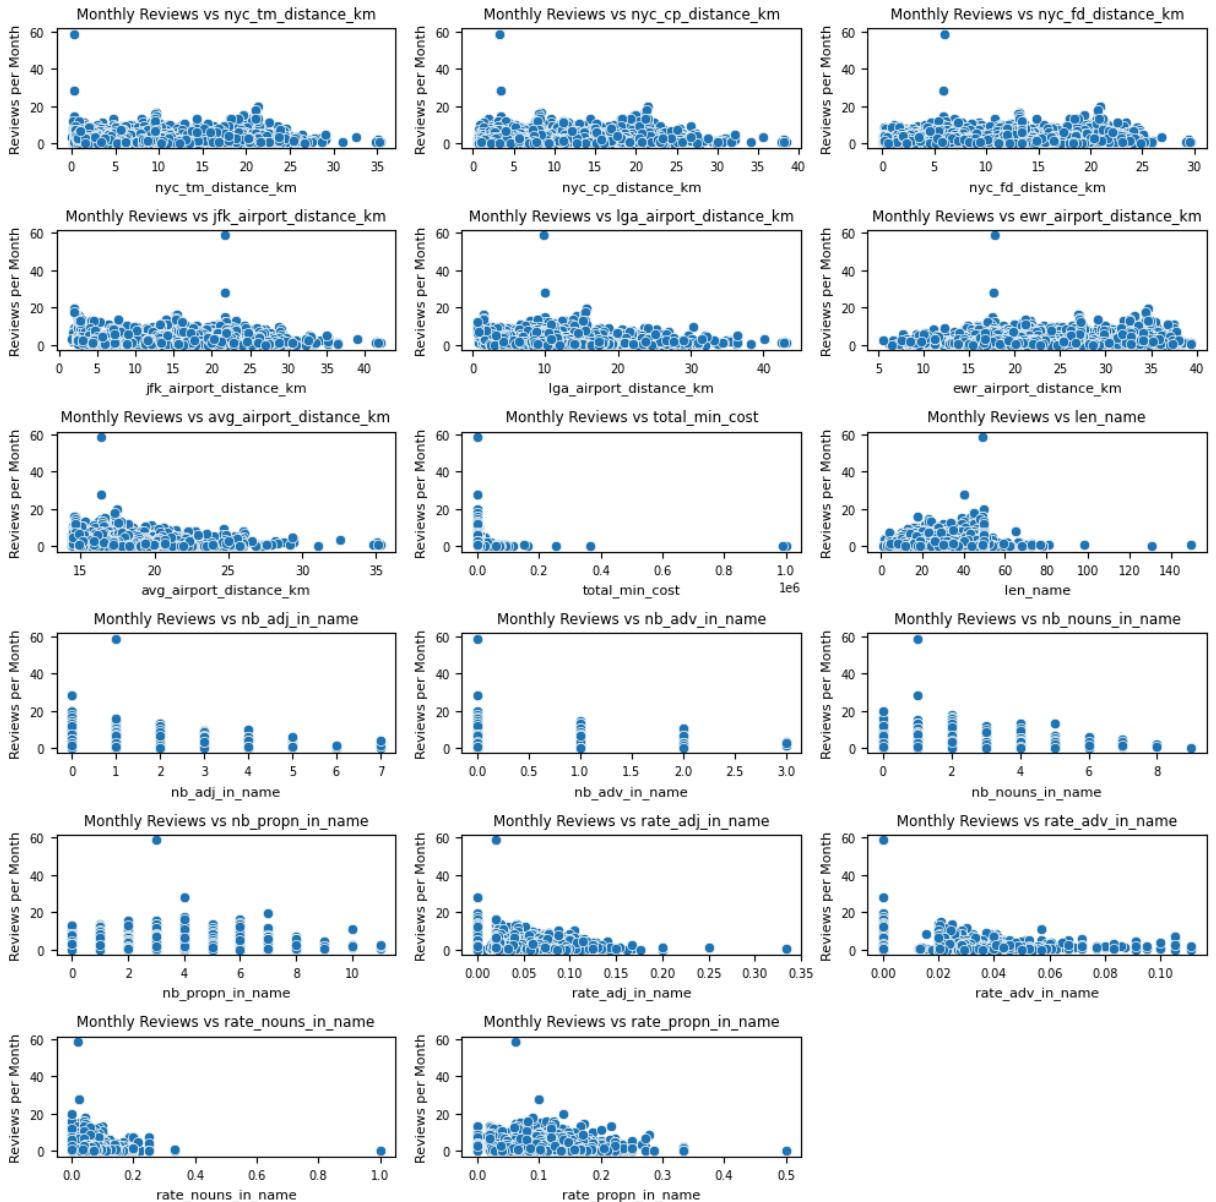
num_cols_per_row = 3
num_rows = int(np.ceil(len(new_features) / num_cols_per_row))

# Create the figure and subplots
fig, axes = plt.subplots(num_rows, num_cols_per_row, figsize=(10, 10))
axes = axes.flatten() # Flatten the 2D array of axes for easier iteration

# Plot histograms for each column
for i, col in enumerate(new_features):
    sns.scatterplot(
        x=df_train[col],
        y=df_train[TARGET],
        ax=axes[i],
    )
    axes[i].set_title(f"Monthly Reviews vs {col}", fontsize=8.5)
    axes[i].set_xlabel(col)
    axes[i].set_ylabel("Reviews per Month")

# Remove any unused subplots if the number of columns is not a multiple of 3
for j in range(len(new_features), len(axes)):
    fig.delaxes(axes[j])

plt.tight_layout()
plt.show()
```



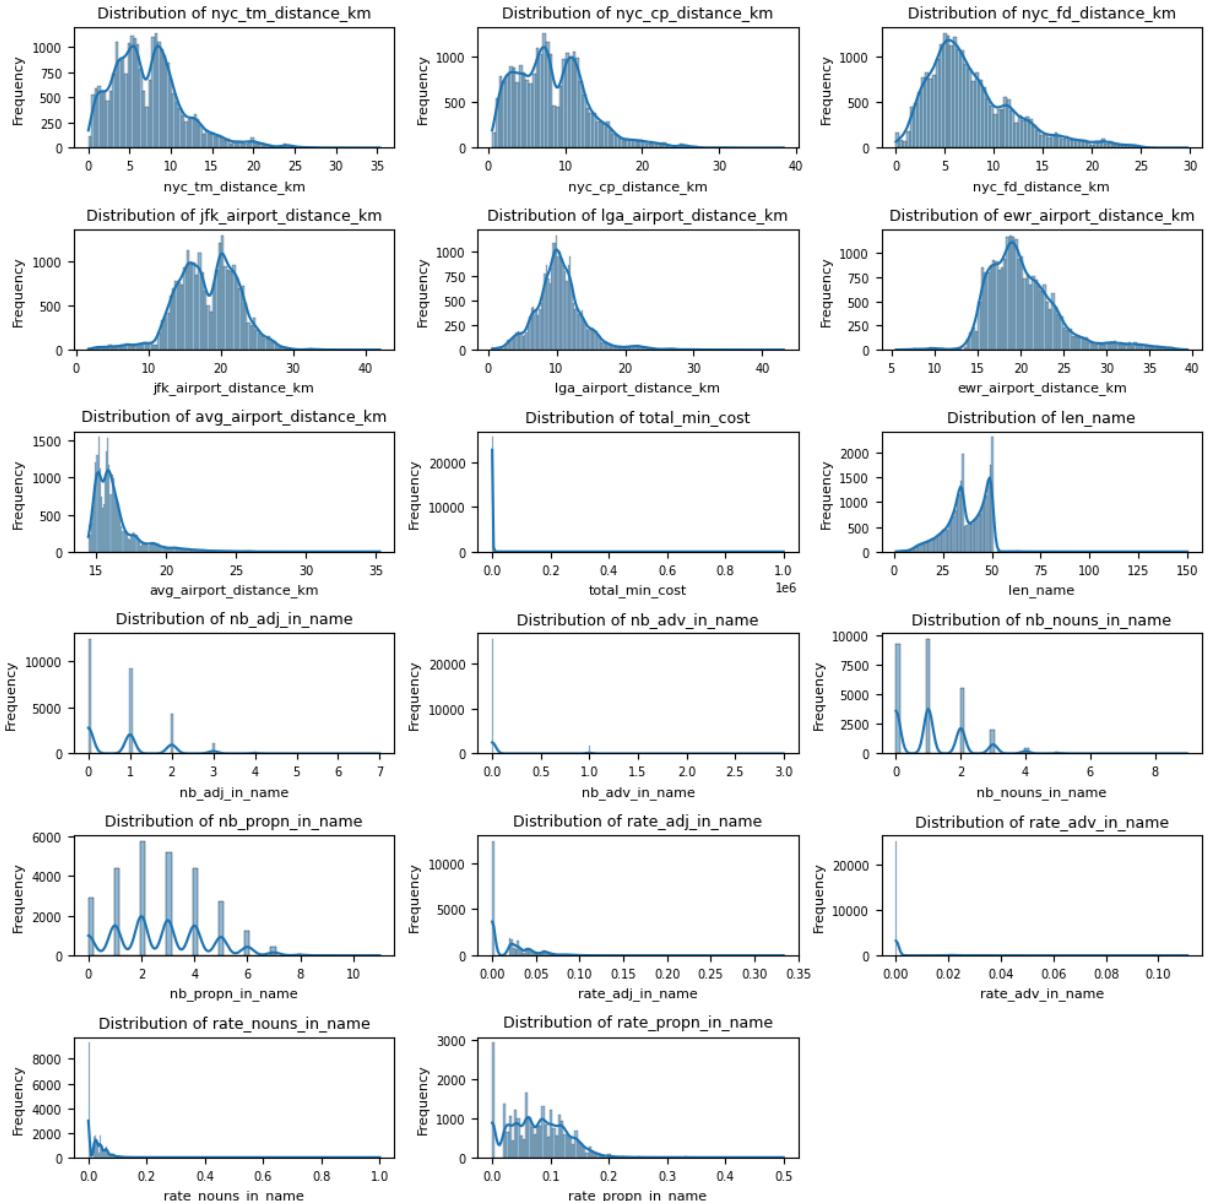
```
In [25]: num_cols_per_row = 3
num_rows = int(np.ceil(len(new_features) / num_cols_per_row))

# Create the figure and subplots
fig, axes = plt.subplots(num_rows, num_cols_per_row, figsize=(10, 10))
axes = axes.flatten() # Flatten the 2D array of axes for easier iteration

# Plot histograms for each column
for i, col in enumerate(new_features):
    sns.histplot(df_train[col], kde=True, ax=axes[i])
    axes[i].set_title(f"Distribution of {col}")
    axes[i].set_xlabel(col)
    axes[i].set_ylabel("Frequency")

# Remove any unused subplots if the number of columns is not a multiple of 3
for j in range(len(new_features), len(axes)):
    fig.delaxes(axes[j])
```

```
plt.tight_layout()
plt.show()
```



```
In [26]: plt.rcParams.update({
    "font.size": 8,
    "axes.titlesize": 9,
    "axes.labelsize": 8,
    "xtick.labelsize": 7,
    "ytick.labelsize": 7,
})

num_cols = df_train.select_dtypes(include="number").columns.tolist()
cont_cols = [c for c in num_cols if c != TARGET]

method = "pearson"

corr_df = df_train[cont_cols + [TARGET]].corr(method=method)

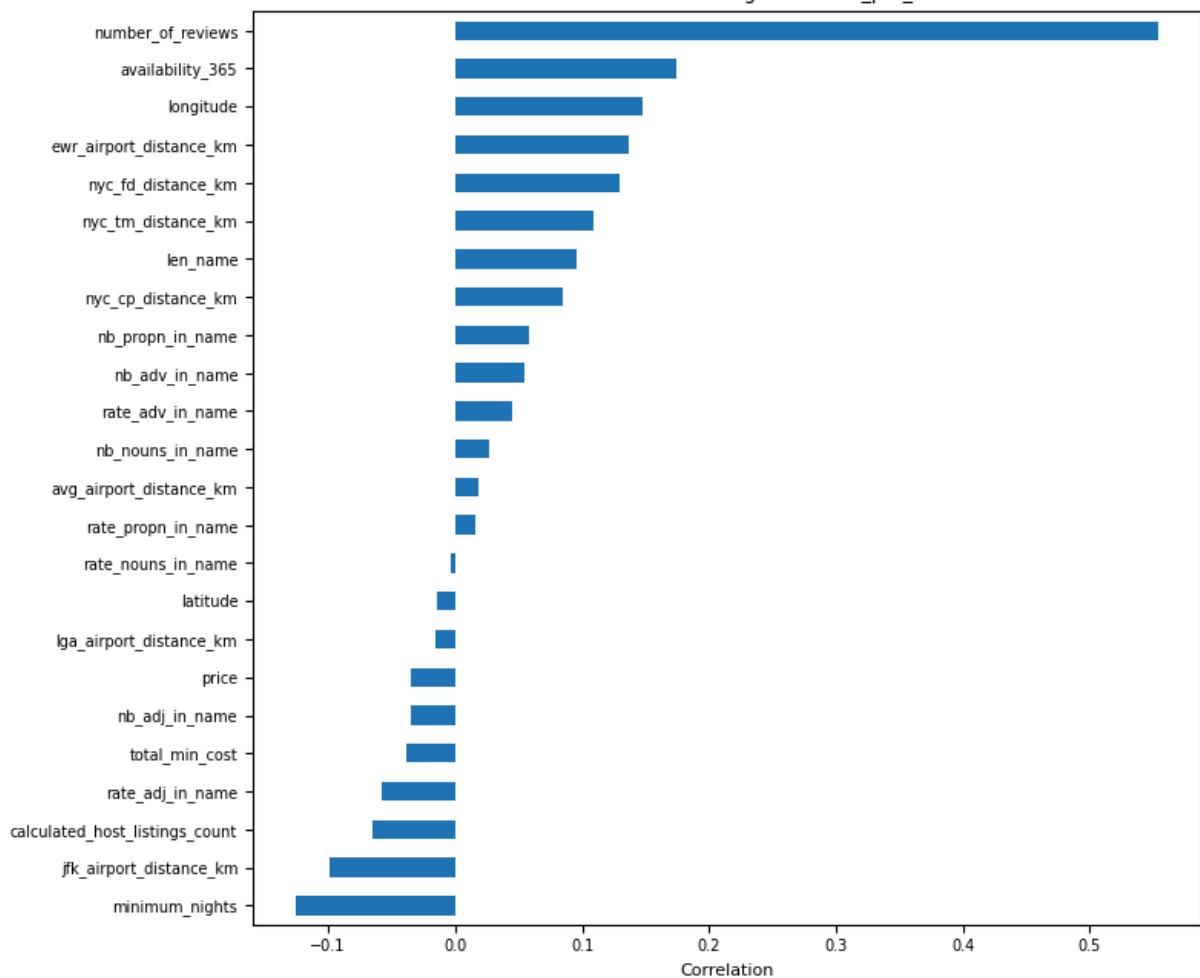
feat_corr = corr_df[TARGET].drop(TARGET).sort_values()
```

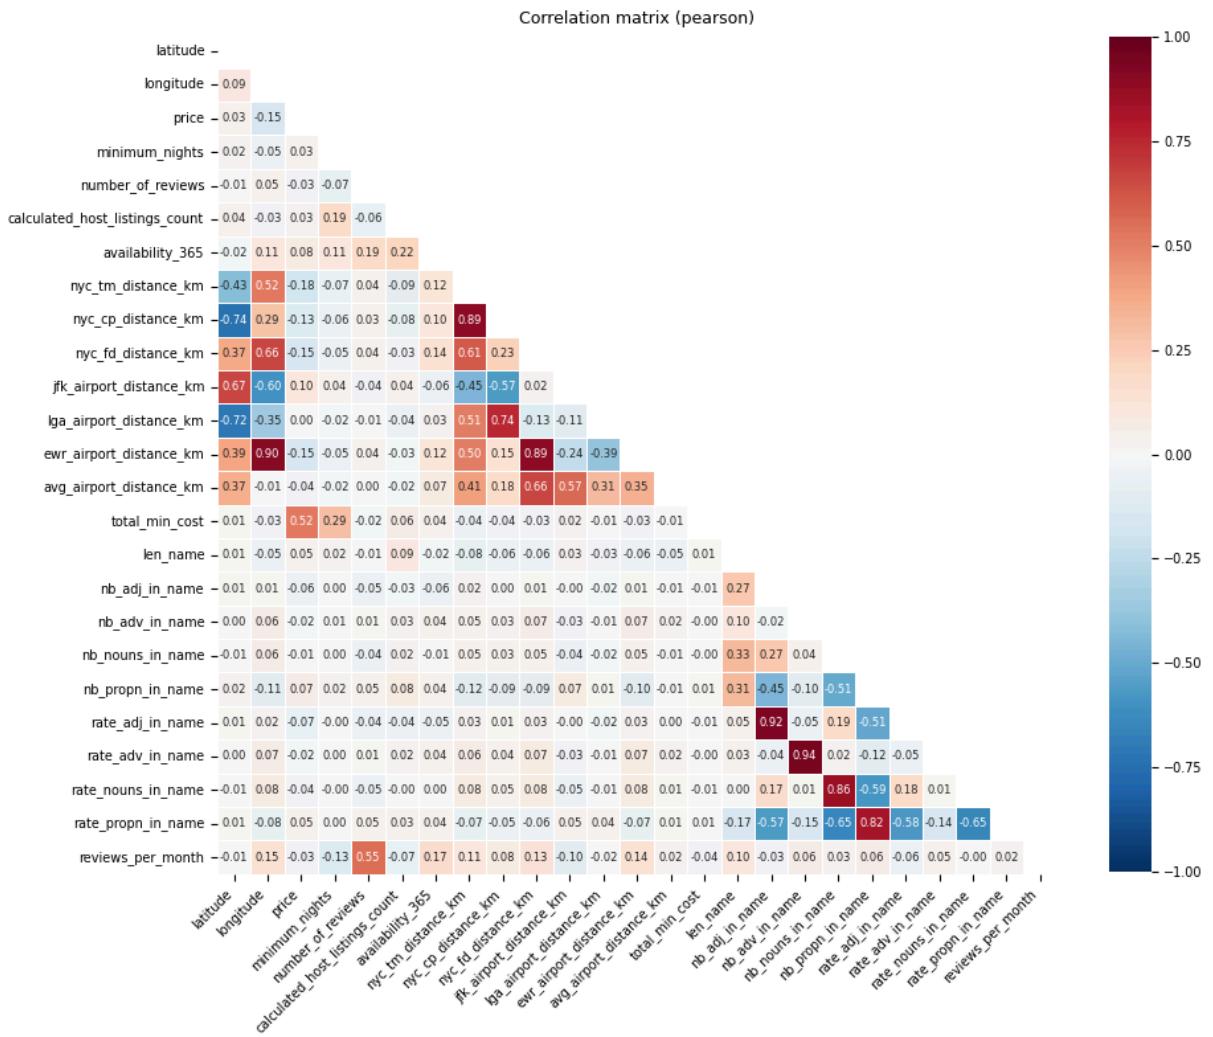
```
plt.figure(figsize=(8, max(4, 0.28 * len(feat_corr))))
feat_corr.plot(kind="barh")
plt.title(f"{{method.title()}} correlation with target: {{TARGET}}", fontsize=9)
plt.xlabel("Correlation", fontsize=8)
plt.yticks(fontsize=7)
plt.xticks(fontsize=7)
plt.tight_layout()
plt.show()

mask = np.triu(np.ones_like(corr_df, dtype=bool))

plt.figure(figsize=(1 + 0.35 * len(corr_df.columns), 1 + 0.35 * len(corr_df.columns))
sns.heatmap(
    corr_df,
    mask=mask,
    center=0,
    vmin=-1, vmax=1,
    cmap="RdBu_r",
    square=True,
    linewidths=0.5,
    annot=True,
    fmt=".2f",
    annot_kws={"size": 6},
    cbar_kws={"shrink": 0.7}
)
plt.title(f"Correlation matrix ({method})", fontsize=9)
plt.xticks(fontsize=7, rotation=45, ha="right")
plt.yticks(fontsize=7)
plt.tight_layout()
plt.show()
```

Pearson correlation with target: reviews\_per\_month





We performed feature engineering to improve the predictive power of the model. The new features include:

- Distance Features: We calculated distances (in kilometers) from each listing to key NYC landmarks and transportation hubs, including Times Square, Central Park, the Financial District, and three major airports (JFK, LaGuardia, Newark). We also computed the average distance to all airports. Proximity to tourist attractions and transportation hubs is likely to affect booking frequency, which impacts the number of reviews received.
- Total Minimum Cost: This feature captures the actual financial commitment required to book a listing by multiplying price by minimum\_nights. This distinguishes between, for example, a low-cost listing with a longer day minimum stay versus an expensive listing with a one day minimum, providing a more realistic measure of booking accessibility.
- Text Features (NLP from Listing Names): We extracted linguistic features from listing names using spaCy, including the length of the name, counts and proportions of different parts of speech (adjectives, adverbs, nouns, proper nouns). Listing names with more descriptive language ("cozy", "spacious") or specific location mentions (proper nouns like "Brooklyn", "Midtown") may attract more bookings and by consequence more reviews.

## 5. Preprocessing and transformations

rubric={accuracy,reasoning}

### Your tasks:

1. Identify different feature types and the transformations you would apply on each feature type.
2. Define a column transformer, if necessary.

Points: 4

```
In [27]: cols_w_outliers = [  
    "price",  
    "minimum_nights",  
    "reviews_per_month",  
    "calculated_host_listings_count",  
]  
df_train[cols_w_outliers].describe(percentiles=[0.5,0.75,0.85,0.90,0.95,0.97,0.99,0
```

```
Out[27]:
```

	price	minimum_nights	reviews_per_month	calculated_host_listings_count
<b>count</b>	27190.000000	27190.000000	27190.000000	27190.000000
<b>mean</b>	142.489702	5.790953	1.368443	3.478301
<b>std</b>	206.368954	16.108038	1.706042	12.036135
<b>min</b>	0.000000	1.000000	0.010000	1.000000
<b>50%</b>	100.000000	2.000000	0.700000	1.000000
<b>75%</b>	169.000000	4.000000	2.000000	2.000000
<b>85%</b>	200.000000	7.000000	3.000000	3.000000
<b>90%</b>	250.000000	14.000000	3.630000	5.000000
<b>95%</b>	330.000000	30.000000	4.675500	9.000000
<b>97%</b>	400.000000	30.000000	5.420000	17.000000
<b>99%</b>	699.000000	31.000000	7.301100	52.000000
<b>99.5%</b>	900.000000	55.165000	8.560550	96.000000
<b>max</b>	10000.000000	999.000000	58.500000	232.000000

```
In [28]: price_outlier_cond = df_train["price"] < np.percentile(df_train["price"], 99)  
minimum_nights_outlier_cond = df_train["minimum_nights"] <= np.percentile(  
    df_train["minimum_nights"], 99  
)
```

```

reviews_per_month_outlier_cond = df_train["reviews_per_month"] <= np.percentile(
    df_train["reviews_per_month"], 99.5
)
listing_count_outlier_cond = df_train["calculated_host_listings_count"] < np.percentile(
    df_train["calculated_host_listings_count"], 99
)

```

In [29]:

```

df_train = df_train.loc[
    price_outlier_cond &
    minimum_nights_outlier_cond &
    reviews_per_month_outlier_cond &
    listing_count_outlier_cond
].copy()

```

In [30]:

```
df_train[cols_w_outliers].describe(percentiles=[0.5,0.75,0.85,0.90,0.95,0.97,0.99,0.995])
```

Out[30]:

	price	minimum_nights	reviews_per_month	calculated_host_listings_count
<b>count</b>	26311.000000	26311.000000	26311.000000	26311.000000
<b>mean</b>	130.162442	4.781992	1.342145	2.424575
<b>std</b>	91.830858	7.152393	1.535088	4.568263
<b>min</b>	0.000000	1.000000	0.010000	1.000000
<b>50%</b>	100.000000	2.000000	0.720000	1.000000
<b>75%</b>	165.000000	4.000000	2.010000	2.000000
<b>85%</b>	200.000000	6.000000	2.990000	3.000000
<b>90%</b>	250.000000	9.000000	3.600000	4.000000
<b>95%</b>	300.000000	30.000000	4.570000	7.000000
<b>97%</b>	360.000000	30.000000	5.200000	11.000000
<b>99%</b>	499.000000	30.000000	6.589000	30.000000
<b>99.5%</b>	550.000000	30.000000	7.330000	34.000000
<b>max</b>	695.000000	31.000000	8.560000	50.000000

In [31]:

```

features = [
    'neighbourhood_group',
    'latitude',
    'longitude',
    'room_type',
    'price',
    'minimum_nights',
    'calculated_host_listings_count',
    'availability_365',
    'nyc_tm_distance_km',
    'nyc_cp_distance_km',
    'nyc_fd_distance_km',
    'jfk_airport_distance_km',
]

```

```

'lga_airport_distance_km',
'ewr_airport_distance_km',
'avg_airport_distance_km',
'total_min_cost',
'len_name',
'nb_adj_in_name',
'nb_adv_in_name',
'nb_nouns_in_name',
'nb_propn_in_name',
'rate_adj_in_name',
'rate_adv_in_name',
'rate_nouns_in_name',
'rate_propn_in_name'
]
categorical_features = [
    'neighbourhood_group',
    'room_type',
]
cont_features = [
    f for f in features if f not in categorical_features
]

```

In [32]: `df_train.to_csv("data/preprocessed/df_train.csv", index=False)`  
`df_test.to_csv("data/preprocessed/df_test.csv", index=False)`

In [33]: `general_preprocessor = make_column_transformer(
 (PolynomialFeatures(degree=2), cont_features),
 (StandardScaler(), cont_features),
 (OneHotEncoder(drop='first', handle_unknown="ignore", sparse_output=False), cat
)
general_preprocessor.set_output(transform="pandas")
general_preprocessor`

Out[33]:

```

graph TD
    CT[ColumnTransformer] --> PF[polynomialfeatures]
    CT --> SS[standardscaler]
    CT --> OH[onehotencoder]
    PF --> P[PolynomialFeatures]
    SS --> S[StandardScaler]
    OH --> O[OneHotEncoder]

```

In [34]: `X_train = df_train[features].copy()`  
`y_train = df_train[TARGET].copy()`  
`X_test = df_test[features].copy()`  
`y_test = df_test[TARGET].copy()`  
  
`X_train_s = general_preprocessor.fit_transform(X_train)`  
`X_test_s = general_preprocessor.transform(X_test)`

### Preprocessing Steps

- a) Outlier Removal: We removed extreme outliers from the training set by capping features at (99 or 99.5th percentile): price, minimum\_nights, reviews\_per\_month, and calculated\_host\_listing\_count.

- b) Feature Scaling: We applied "StandardScaler" to all continuous features to ensure features are on comparable scales.
- c) Categorical Encoding: Applied "OneHotEncoder" with drop='first' to try to avoid multicollinearity and handle\_unknown='ignore' for robustness to new categories in test data.
- d) Final Feature Set: 25 features total (23 continuous, 2 categorical → ~26 after one-hot encoding).

## 6. Baseline model

`rubric={accuracy}`

**Your tasks:**

1. Train a baseline model for your task and report its performance.

*Points:* 2

```
In [35]: from sklearn.dummy import DummyRegressor

dummy_regr = DummyRegressor(strategy="mean")
dummy_regr.fit(X_train_s, y_train)

y_train_pred = dummy_regr.predict(X_train_s)
y_test_pred = dummy_regr.predict(X_test_s)

print("BASELINE MODEL RESULTS:")
print("TRAIN MSE", round(mean_squared_error(y_true=y_train, y_pred=y_train_pred), 3))
print("TEST MSE", round(mean_squared_error(y_true=y_test, y_pred=y_test_pred), 3))

print("TRAIN R2", round(r2_score(y_true=y_train, y_pred=y_train_pred), 3))
print("TEST R2", round(r2_score(y_true=y_test, y_pred=y_test_pred), 3))

print(
    "TRAIN RMSE", round(root_mean_squared_error(y_true=y_train, y_pred=y_train_pred)
)
print("TEST RMSE", round(root_mean_squared_error(y_true=y_test, y_pred=y_test_pred))
```

BASELINE MODEL RESULTS:

TRAIN MSE 2.356

TEST MSE 2.623

TRAIN R2 0.0

TEST R2 -0.001

TRAIN RMSE 1.535

TEST RMSE 1.62

We chose a baseline using DummyRegressor with a mean prediction strategy, which predicts the training set mean for all observations.

This baseline represents the simplest possible model and provides a benchmark for evaluating more sophisticated approaches. Any model with R squared > 0 demonstrates improvement over this naive baseline. The RMSE of approximately 1.6 represents the standard deviation of the target variable, indicating that on average, predictions are off by about 1.6 reviews per month when simply predicting the mean.

## 7. Linear models

`rubric={accuracy,reasoning}`

### Your tasks:

1. Try a linear model as a first real attempt.
2. Carry out hyperparameter tuning to explore different values for the regularization hyperparameter.
3. Report cross-validation scores along with standard deviation.
4. Summarize your results.

*Points:* 8

Our Ridge regression model is performing not well with R squared = 0.253, meaning only 25% of variance is explained. Your interpretation is on track - there's likely significant omitted variable bias (OVB).

The main problems that we see are:

a) R squared: Test determination coefficient of 0.253 is indeed weak. The model is missing some important predictors. b) Also, there is a high alpha (266,322): This extreme regularization is shrinking coefficients heavily, suggesting multicollinearity and that our features aren't predictive. c) We have very tiny coefficients and as are scaled this means features have minimal impact.

Interpretation of Top Polynomial Features:

- Minimum Nights\*\*2 (0.0054): A one-unit increase in minimum\_nights to the square increases the target by 0.0054. The squared term suggests a non linear relationship that is extremely long minimum stay requirements might have diminishing or accelerating effects on your target variable.
- Nyc\_fd\_distance\_km \* Jfk\_airport\_distance\_km (0.0048): This interaction suggests that being far from both Fire Department landmarks and JFK airport simultaneously has a small positive effect. This could capture outer borough or remote listings.

In [41]: `from sklearn.linear_model import Ridge`

```
    raw_ridge = make_pipeline(  
        general_preprocessor,  
        Ridge()  
)
```

```
In [58]: from sklearn.model_selection import RandomizedSearchCV  
from scipy.stats import loguniform, uniform, randint  
  
param_grid = {  
    "ridge_alpha": loguniform(1e-10, 1e10),  
}  
  
random_search_ridge = RandomizedSearchCV(  
    raw_ridge,  
    param_distributions=param_grid,  
    n_iter=100,  
    n_jobs=-1,  
    random_state=RANDOM_STATE,  
    return_train_score=True,  
)  
  
random_search_ridge.fit(X_train, y_train)  
  
print(  
    "Random Search best model score: \t %0.3f" % random_search_ridge.best_score_  
)  
print(  
    "Random Search best alpha: \t\t\t %0.3f" %  
    random_search_ridge.best_params_["ridge_alpha"]  
)  
  
pd.DataFrame(random_search_ridge.cv_results_)[  
    [  
        "mean_test_score",  
        "param_ridge_alpha",  
        "mean_fit_time",  
        "rank_test_score",  
    ]  
].set_index("rank_test_score").sort_index()
```

```
Random Search best model score:          0.253  
Random Search best alpha:                266321.933
```

Out[58]:

	mean_test_score	param_ridge_alpha	mean_fit_time
rank_test_score			
1	0.252686	2.663219e+05	1.300377
2	0.252686	2.626296e+05	1.206110
3	0.252686	2.785450e+05	1.187796
4	0.252680	3.181680e+05	1.190700
5	0.252632	1.643240e+05	1.406712
...	...	...	...
96	0.241700	1.033442e+09	1.023568
97	0.240047	2.054190e+09	1.244403
98	0.239541	2.464294e+09	1.213482
99	0.239498	2.501480e+09	1.114498
100	0.236836	5.466870e+09	1.231342

100 rows × 3 columns

In [53]:

```
from sklearn.model_selection import cross_validate

def mean_std_cross_val_scores(model, X_train, y_train, scoring, **kwargs):
    scores = cross_validate(
        model,
        X_train,
        y_train,
        return_train_score=True,
        scoring=scoring,
        **kwargs
    )
    df = pd.DataFrame(scores)
    return df.mean(), df.std()

score_types_reg = {
    "neg_root_mean_squared_error": "neg_root_mean_squared_error",
    "r2": "r2",
}

results_mean, results_std = {}, {}

mean, std = mean_std_cross_val_scores(
    random_search_ridge.best_estimator_, X_train, y_train, scoring=score_types_reg
)
results_mean["Ridge"] = mean
results_std["Ridge"] = std

mean_df = pd.DataFrame(results_mean).T
```

```
std_df = pd.DataFrame(results_std).T
results_df = mean_df.round(3).astype(str) + " ± " + std_df.round(3).astype(str)
results_df.T
```

Out[53]:

Ridge	
<b>fit_time</b>	0.183 ± 0.028
<b>score_time</b>	0.037 ± 0.007
<b>test_neg_root_mean_squared_error</b>	-1.327 ± 0.01
<b>train_neg_root_mean_squared_error</b>	-1.314 ± 0.003
<b>test_r2</b>	0.253 ± 0.011
<b>train_r2</b>	0.268 ± 0.002

In [72]:

```
import mglearn

best_pipe = random_search_ridge.best_estimator_

ridge = best_pipe.named_steps["ridge"]
coef = np.ravel(ridge.coef_)

feature_names = best_pipe[:-1].get_feature_names_out()

coef_df = pd.DataFrame({"feature": feature_names, "coefficient": coef})
coef_df["magnitude"] = coef_df["coefficient"].abs()

split = coef_df["feature"].str.split("__", n=1, expand=True)

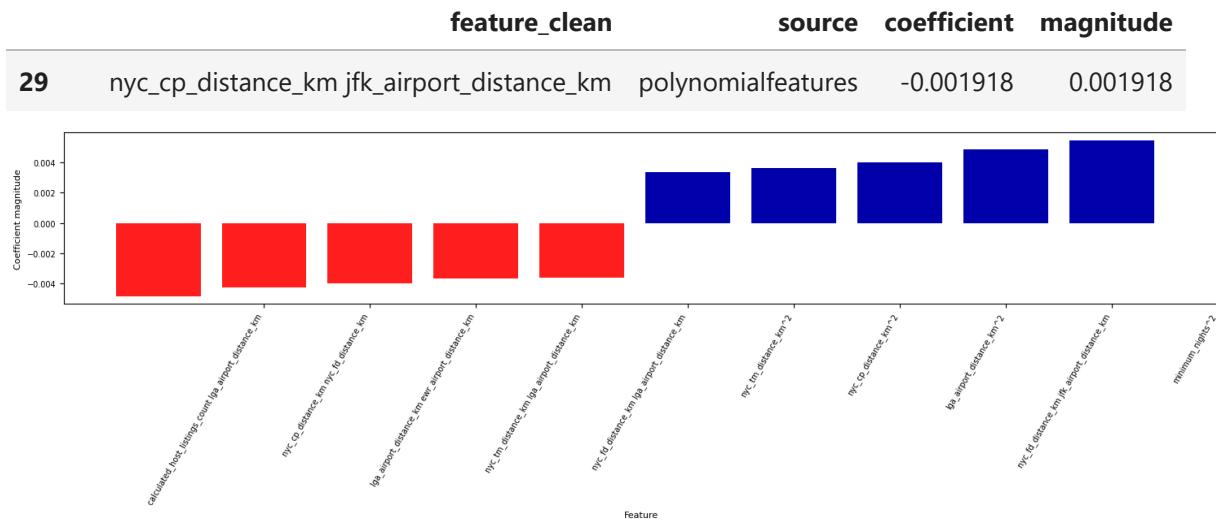
if split.shape[1] == 2:
    coef_df["source"] = split[0]
    coef_df["feature_clean"] = split[1]
else:
    coef_df["source"] = None
    coef_df["feature_clean"] = split[0]

top30 = (coef_df.sort_values("magnitude", ascending=False)
          .head(30)
          .reset_index(drop=True))

display(top30[["feature_clean", "source", "coefficient", "magnitude"]])

mglearn.tools.visualize_coefficients(
    coef,
    coef_df["feature_clean"].to_numpy(),
    n_top_features=5
)
plt.tight_layout()
plt.show()
```

		feature_clean	source	coefficient	magnitude
0		minimum_nights^2	polynomialfeatures	0.005445	0.005445
1	nyc_fd_distance_km jfk_airport_distance_km		polynomialfeatures	0.004828	0.004828
2	calculated_host_listings_count lga_airport_dis...		polynomialfeatures	-0.004824	0.004824
3	nyc_cp_distance_km nyc_fd_distance_km		polynomialfeatures	-0.004268	0.004268
4	lga_airport_distance_km ewr_airport_distance_km		polynomialfeatures	-0.003994	0.003994
5	lga_airport_distance_km^2		polynomialfeatures	0.003979	0.003979
6	nyc_tm_distance_km lga_airport_distance_km		polynomialfeatures	-0.003644	0.003644
7	nyc_cp_distance_km^2		polynomialfeatures	0.003632	0.003632
8	nyc_fd_distance_km lga_airport_distance_km		polynomialfeatures	-0.003599	0.003599
9	nyc_tm_distance_km^2		polynomialfeatures	0.003341	0.003341
10	jfk_airport_distance_km^2		polynomialfeatures	0.003245	0.003245
11	nyc_fd_distance_km nb_adv_in_name		polynomialfeatures	-0.003179	0.003179
12	calculated_host_listings_count nb_adj_in_name		polynomialfeatures	0.002957	0.002957
13	len_name nb_propn_in_name		polynomialfeatures	0.002812	0.002812
14	nyc_fd_distance_km^2		polynomialfeatures	0.002794	0.002794
15	calculated_host_listings_count nyc_fd_distance_km		polynomialfeatures	-0.002732	0.002732
16	calculated_host_listings_count jfk_airport_dis...		polynomialfeatures	0.002645	0.002645
17	nyc_cp_distance_km ewr_airport_distance_km		polynomialfeatures	-0.002590	0.002590
18	price rate_propn_in_name		polynomialfeatures	0.002547	0.002547
19	minimum_nights nyc_tm_distance_km		polynomialfeatures	0.002510	0.002510
20	len_name nb_adj_in_name		polynomialfeatures	0.002282	0.002282
21	room_type_Private room	onehotencoder		-0.002277	0.002277
22	nyc_tm_distance_km jfk_airport_distance_km		polynomialfeatures	-0.002241	0.002241
23	jfk_airport_distance_km nb_adj_in_name		polynomialfeatures	0.002224	0.002224
24	minimum_nights calculated_host_listings_count		polynomialfeatures	-0.002194	0.002194
25	minimum_nights nyc_fd_distance_km		polynomialfeatures	-0.002100	0.002100
26	nyc_tm_distance_km nb_adj_in_name		polynomialfeatures	-0.001986	0.001986
27	longitude jfk_airport_distance_km		polynomialfeatures	0.001956	0.001956
28	availability_365 rate_propn_in_name		polynomialfeatures	0.001950	0.001950



## 8. Different models

rubric={accuracy,reasoning}

### Your tasks:

1. Try out three other models aside from the linear model.
2. Summarize your results in terms of overfitting/underfitting and fit and score times.  
Can you beat the performance of the linear model?

Points: 10

Now, we decide to compare three models in order to improve the predictive performance and we have the following:

1. The best model performance (by the R squared on test set) is Gradient Boosting (GBR) with 0.369, then Random Forest (0.338), Ridge (0.253) and Decision Tree is the worst (0.30).

So, we can say that GBR improves the previous Ridge because it can explains 37% of the variance of the model and also has the lowest prediction error (RMSE of 1.219). Also, it has overfitting but not so much (0.563 in train vs 0.369 on test). Random Forest is also close but it has a great overfit (0.847 in train but in test 0.338).

In summary, with that we can be certain that Ridge is not so great because it underfits and needs a lot of regularization. Decision Tree is the worst model of all (performing worse than just predicting the mean).

In [57]: `from sklearn.ensemble import HistGradientBoostingRegressor`

```
models = {
```

```

    "DecisionTree": DecisionTreeRegressor(
        random_state=RANDOM_STATE,
        max_depth=200,
        max_features=100
    ),
    "RandomForest": RandomForestRegressor(
        random_state=RANDOM_STATE,
        max_depth=20,
        n_estimators=50,
        n_jobs=-1,
        max_features=100,
    ),
    "GBR": HistGradientBoostingRegressor(
        max_depth=8,
        learning_rate=0.05,
        max_iter=300,
        random_state=RANDOM_STATE
    )
}

for name, model in models.items():
    pipe = make_pipeline(
        general_preprocessor,
        model
    )
    mean, std = mean_std_cross_val_scores(
        pipe, X_train, y_train, scoring(score_types_reg)
    )
    results_mean[name] = mean
    results_std[name] = std
    print(name)

mean_df = pd.DataFrame(results_mean).T
std_df = pd.DataFrame(results_std).T
results_df = mean_df.round(3).astype(str) + " ± " + std_df.round(3).astype(str)
results_df.T

```

DecisionTree  
 RandomForest  
 GBR

Out[57]:

		Ridge	RandomForest	DecisionTree	GBR
	<b>fit_time</b>	0.183 ± 0.028	12.006 ± 0.815	3.04 ± 0.09	8.828 ± 0.867
	<b>score_time</b>	0.037 ± 0.007	0.077 ± 0.01	0.038 ± 0.011	0.109 ± 0.021
	<b>test_neg_root_mean_squared_error</b>	-1.327 ± 0.01	-1.248 ± 0.018	-1.75 ± 0.012	-1.219 ± 0.018
	<b>train_neg_root_mean_squared_error</b>	-1.314 ± 0.003	-0.601 ± 0.011	-0.0 ± 0.0	-1.015 ± 0.025
	<b>test_r2</b>	0.253 ± 0.011	0.338 ± 0.008	-0.3 ± 0.024	0.369 ± 0.007
	<b>train_r2</b>	0.268 ± 0.002	0.847 ± 0.005	1.0 ± 0.0	0.563 ± 0.021

## 9. Feature selection (Challenging)

rubric={reasoning}

**Your tasks:**

Make some attempts to select relevant features. You may try `RFEcv`, forward/backward selection or L1 regularization for this. Do the results improve with feature selection?

Summarize your results. If you see improvements in the results, keep feature selection in your pipeline. If not, you may abandon it in the next exercises unless you think there are other benefits with using fewer features.

Points: 0.5

Type your answer here, replacing this text.

In [14]:

```
X_train_soft = X_train.copy()

soft_features = [
    "neighbourhood_group",
    "latitude",
    "longitude",
    "room_type",
    "price",
    "minimum_nights",
    "calculated_host_listings_count",
    "availability_365",
    "nyc_tm_distance_km",
    "nyc_cp_distance_km",
    "nyc_fd_distance_km",
    "jfk_airport_distance_km",
```

```

    "lga_airport_distance_km",
    "ewr_airport_distance_km",
    "avg_airport_distance_km",
    "total_min_cost",
    "len_name",
    "nb_adj_in_name",
    "nb_adv_in_name",
    "nb_nouns_in_name",
    "nb_propn_in_name",
    "rate_adj_in_name",
    "rate_adv_in_name",
    "rate_nouns_in_name",
    "rate_propn_in_name",
]
soft_cont_features = [f for f in soft_features if f not in categorical_features]
soft_preprocessor = make_column_transformer(
    (PolynomialFeatures(degree=2), cont_features),
    (StandardScaler(), cont_features),
    (
        OneHotEncoder(drop="first", handle_unknown="ignore", sparse_output=False),
        categorical_features,
    ),
)
soft_preprocessor.set_output(transform="pandas")
X_train_soft_s = soft_preprocessor.fit_transform(X_train_soft)
X_train_soft_s["__RANDOM_1__"] = np.random.rand(len(X_train_soft))
X_train_soft_s["__RANDOM_2__"] = np.random.rand(len(X_train_soft))
soft_rf = RandomForestRegressor(n_jobs=-1, random_state=RANDOM_STATE, max_depth=6)
soft_rf.fit(X_train_soft_s, y_train)

```

Out[14]:

▼ RandomForestRegressor 
  
► Parameters

In [15]:

```
explainer = shap.TreeExplainer(soft_rf)
shap_values = explainer.shap_values(X_train_soft_s)
```

In [32]:

```
random_feats = ["__RANDOM_1__", "__RANDOM_2__"]
returned_features = soft_preprocessor.get_feature_names_out()
clean_features = [
    v.split("__", 1)[1] for v in returned_features if v not in random_feats
] + random_feats
ranked_features, selected_features = get_important_features_important_than(
    shap_values, clean_features, random_feats
)
print(f"{len(selected_features)} valuable features selected")
selected_features
```

90 valuable features selected

```
Out[32]: ['latitude availability_365',
 'minimum_nights avg_airport_distance_km',
 'price availability_365',
 'longitude minimum_nights',
 'availability_365 ewr_airport_distance_km',
 'longitude availability_365',
 'latitude minimum_nights',
 'calculated_host_listings_count len_name',
 'availability_365 jfk_airport_distance_km',
 'latitude len_name',
 'calculated_host_listings_count ewr_airport_distance_km',
 'availability_365 nyc_fd_distance_km',
 'longitude len_name',
 'minimum_nights jfk_airport_distance_km',
 'room_type_Private room',
 'calculated_host_listings_count nyc_fd_distance_km',
 'minimum_nights calculated_host_listings_count',
 'avg_airport_distance_km len_name',
 'availability_365 len_name',
 'minimum_nights total_min_cost',
 'availability_365 avg_airport_distance_km',
 'availability_365 nb_propn_in_name',
 'nyc_tm_distance_km lga_airport_distance_km',
 'minimum_nights ewr_airport_distance_km',
 'minimum_nights availability_365',
 'latitude calculated_host_listings_count',
 'availability_365 rate_adj_in_name',
 'minimum_nights lga_airport_distance_km',
 'nyc_tm_distance_km total_min_cost',
 'minimum_nights nyc_tm_distance_km',
 'calculated_host_listings_count availability_365',
 'longitude calculated_host_listings_count',
 'jfk_airport_distance_km len_name',
 'ewr_airport_distance_km len_name',
 'nyc_tm_distance_km nyc_cp_distance_km',
 'calculated_host_listings_count total_min_cost',
 'price calculated_host_listings_count',
 'availability_365 rate_propn_in_name',
 'nyc_tm_distance_km ewr_airport_distance_km',
 'jfk_airport_distance_km ewr_airport_distance_km',
 'price nyc_tm_distance_km',
 'minimum_nights len_name',
 'ewr_airport_distance_km total_min_cost',
 'total_min_cost len_name',
 'nyc_cp_distance_km total_min_cost',
 'len_name nb_propn_in_name',
 'availability_365 rate_nouns_in_name',
 'nyc_fd_distance_km lga_airport_distance_km',
 'avg_airport_distance_km total_min_cost',
 'minimum_nights rate_adj_in_name',
 'lga_airport_distance_km total_min_cost',
 'nyc_tm_distance_km nyc_fd_distance_km',
 'calculated_host_listings_count avg_airport_distance_km',
 'longitude total_min_cost',
 'nyc_fd_distance_km jfk_airport_distance_km',
 'minimum_nights nyc_fd_distance_km',
```

```
'availability_365 nyc_tm_distance_km',
'rate_adv_in_name rate_propn_in_name',
'jfk_airport_distance_km lga_airport_distance_km',
'jfk_airport_distance_km avg_airport_distance_km',
'price ewr_airport_distance_km',
'calculated_host_listings_count jfk_airport_distance_km',
'jfk_airport_distance_km total_min_cost',
'nyc_tm_distance_km len_name',
'price avg_airport_distance_km',
'ewr_airport_distance_km nb_propn_in_name',
'calculated_host_listings_count nyc_tm_distance_km',
'nyc_fd_distance_km total_min_cost',
'nyc_tm_distance_km nb_propn_in_name',
'nyc_tm_distance_km jfk_airport_distance_km',
'calculated_host_listings_count rate_propn_in_name',
'lga_airport_distance_km avg_airport_distance_km',
'nyc_fd_distance_km rate_propn_in_name',
'nyc_tm_distance_km rate_adj_in_name',
'nyc_cp_distance_km jfk_airport_distance_km',
'latitude longitude',
'calculated_host_listings_count lga_airport_distance_km',
'ewr_airport_distance_km rate_propn_in_name',
'price total_min_cost',
'calculated_host_listings_count nyc_cp_distance_km',
'nb_nouns_in_name nb_propn_in_name',
'availability_365 total_min_cost',
'nyc_fd_distance_km nb_propn_in_name',
'latitude avg_airport_distance_km',
'longitude nyc_tm_distance_km',
'nyc_tm_distance_km rate_propn_in_name',
'lga_airport_distance_km len_name',
'nyc_cp_distance_km rate_propn_in_name',
'longitude jfk_airport_distance_km',
'nyc_cp_distance_km nb_adj_in_name']
```

```
In [26]: for idx, v in enumerate(ranked_features):
    print(f"{idx + 1} - {v[0]} -> shap: {round(v[1],4)}")
```

1 - latitude availability\_365 -> shap: 0.2634  
2 - minimum\_nights avg\_airport\_distance\_km -> shap: 0.2056  
3 - price availability\_365 -> shap: 0.1541  
4 - longitude minimum\_nights -> shap: 0.1267  
5 - availability\_365 ewr\_airport\_distance\_km -> shap: 0.0821  
6 - longitude availability\_365 -> shap: 0.0648  
7 - latitude minimum\_nights -> shap: 0.0424  
8 - calculated\_host\_listings\_count len\_name -> shap: 0.0307  
9 - availability\_365 jfk\_airport\_distance\_km -> shap: 0.0274  
10 - latitude len\_name -> shap: 0.025  
11 - calculated\_host\_listings\_count ewr\_airport\_distance\_km -> shap: 0.0249  
12 - availability\_365 nyc\_fd\_distance\_km -> shap: 0.0244  
13 - longitude len\_name -> shap: 0.0206  
14 - minimum\_nights jfk\_airport\_distance\_km -> shap: 0.0168  
15 - room\_type\_Private room -> shap: 0.0151  
16 - calculated\_host\_listings\_count nyc\_fd\_distance\_km -> shap: 0.0111  
17 - minimum\_nights calculated\_host\_listings\_count -> shap: 0.0096  
18 - avg\_airport\_distance\_km len\_name -> shap: 0.0086  
19 - availability\_365 len\_name -> shap: 0.007  
20 - minimum\_nights total\_min\_cost -> shap: 0.0065  
21 - availability\_365 avg\_airport\_distance\_km -> shap: 0.0063  
22 - availability\_365 nb\_propn\_in\_name -> shap: 0.0057  
23 - nyc\_tm\_distance\_km lga\_airport\_distance\_km -> shap: 0.005  
24 - minimum\_nights ewr\_airport\_distance\_km -> shap: 0.0049  
25 - minimum\_nights availability\_365 -> shap: 0.0045  
26 - latitude calculated\_host\_listings\_count -> shap: 0.0029  
27 - availability\_365 rate\_adj\_in\_name -> shap: 0.0029  
28 - minimum\_nights lga\_airport\_distance\_km -> shap: 0.0027  
29 - nyc\_tm\_distance\_km total\_min\_cost -> shap: 0.0026  
30 - minimum\_nights nyc\_tm\_distance\_km -> shap: 0.0025  
31 - calculated\_host\_listings\_count availability\_365 -> shap: 0.0024  
32 - longitude calculated\_host\_listings\_count -> shap: 0.0024  
33 - jfk\_airport\_distance\_km len\_name -> shap: 0.0023  
34 - ewr\_airport\_distance\_km len\_name -> shap: 0.0022  
35 - nyc\_tm\_distance\_km nyc\_cp\_distance\_km -> shap: 0.0022  
36 - calculated\_host\_listings\_count total\_min\_cost -> shap: 0.0021  
37 - price calculated\_host\_listings\_count -> shap: 0.002  
38 - availability\_365 rate\_propn\_in\_name -> shap: 0.0018  
39 - nyc\_tm\_distance\_km ewr\_airport\_distance\_km -> shap: 0.0018  
40 - jfk\_airport\_distance\_km ewr\_airport\_distance\_km -> shap: 0.0017  
41 - price nyc\_tm\_distance\_km -> shap: 0.0015  
42 - minimum\_nights len\_name -> shap: 0.0015  
43 - ewr\_airport\_distance\_km total\_min\_cost -> shap: 0.0014  
44 - total\_min\_cost len\_name -> shap: 0.0014  
45 - nyc\_cp\_distance\_km total\_min\_cost -> shap: 0.0014  
46 - len\_name nb\_propn\_in\_name -> shap: 0.0013  
47 - availability\_365 rate\_nouns\_in\_name -> shap: 0.0013  
48 - nyc\_fd\_distance\_km lga\_airport\_distance\_km -> shap: 0.0012  
49 - avg\_airport\_distance\_km total\_min\_cost -> shap: 0.0012  
50 - minimum\_nights rate\_adj\_in\_name -> shap: 0.0011  
51 - lga\_airport\_distance\_km total\_min\_cost -> shap: 0.0011  
52 - nyc\_tm\_distance\_km nyc\_fd\_distance\_km -> shap: 0.0011  
53 - calculated\_host\_listings\_count avg\_airport\_distance\_km -> shap: 0.001  
54 - longitude total\_min\_cost -> shap: 0.0009  
55 - nyc\_fd\_distance\_km jfk\_airport\_distance\_km -> shap: 0.0009  
56 - minimum\_nights nyc\_fd\_distance\_km -> shap: 0.0009

57 - availability\_365 nyc\_tm\_distance\_km -> shap: 0.0008  
58 - rate\_adv\_in\_name rate\_propn\_in\_name -> shap: 0.0008  
59 - jfk\_airport\_distance\_km lga\_airport\_distance\_km -> shap: 0.0007  
60 - jfk\_airport\_distance\_km avg\_airport\_distance\_km -> shap: 0.0007  
61 - price ewr\_airport\_distance\_km -> shap: 0.0007  
62 - calculated\_host\_listings\_count jfk\_airport\_distance\_km -> shap: 0.0007  
63 - jfk\_airport\_distance\_km total\_min\_cost -> shap: 0.0007  
64 - nyc\_tm\_distance\_km len\_name -> shap: 0.0007  
65 - price avg\_airport\_distance\_km -> shap: 0.0006  
66 - ewr\_airport\_distance\_km nb\_propn\_in\_name -> shap: 0.0006  
67 - calculated\_host\_listings\_count nyc\_tm\_distance\_km -> shap: 0.0006  
68 - nyc\_fd\_distance\_km total\_min\_cost -> shap: 0.0006  
69 - nyc\_tm\_distance\_km nb\_propn\_in\_name -> shap: 0.0006  
70 - nyc\_tm\_distance\_km jfk\_airport\_distance\_km -> shap: 0.0006  
71 - calculated\_host\_listings\_count rate\_propn\_in\_name -> shap: 0.0006  
72 - lga\_airport\_distance\_km avg\_airport\_distance\_km -> shap: 0.0006  
73 - nyc\_fd\_distance\_km rate\_propn\_in\_name -> shap: 0.0006  
74 - nyc\_tm\_distance\_km rate\_adj\_in\_name -> shap: 0.0006  
75 - nyc\_cp\_distance\_km jfk\_airport\_distance\_km -> shap: 0.0005  
76 - availability\_365 -> shap: 0.0005  
77 - latitude longitude -> shap: 0.0005  
78 - calculated\_host\_listings\_count lga\_airport\_distance\_km -> shap: 0.0005  
79 - ewr\_airport\_distance\_km rate\_propn\_in\_name -> shap: 0.0005  
80 - price total\_min\_cost -> shap: 0.0005  
81 - calculated\_host\_listings\_count nyc\_cp\_distance\_km -> shap: 0.0004  
82 - nb\_nouns\_in\_name nb\_propn\_in\_name -> shap: 0.0004  
83 - availability\_365 total\_min\_cost -> shap: 0.0004  
84 - nyc\_fd\_distance\_km nb\_propn\_in\_name -> shap: 0.0004  
85 - latitude avg\_airport\_distance\_km -> shap: 0.0004  
86 - longitude nyc\_tm\_distance\_km -> shap: 0.0004  
87 - nyc\_tm\_distance\_km rate\_propn\_in\_name -> shap: 0.0004  
88 - lga\_airport\_distance\_km len\_name -> shap: 0.0004  
89 - nyc\_cp\_distance\_km rate\_propn\_in\_name -> shap: 0.0004  
90 - longitude jfk\_airport\_distance\_km -> shap: 0.0004  
91 - nyc\_cp\_distance\_km nb\_adj\_in\_name -> shap: 0.0004  
92 - RANDOM\_2 -> shap: 0.0004  
93 - RANDOM\_1 -> shap: 0.0004  
94 - latitude rate\_adj\_in\_name -> shap: 0.0004  
95 - calculated\_host\_listings\_count nb\_propn\_in\_name -> shap: 0.0004  
96 - availability\_365 lga\_airport\_distance\_km -> shap: 0.0003  
97 - nyc\_cp\_distance\_km nb\_propn\_in\_name -> shap: 0.0003  
98 - nyc\_fd\_distance\_km len\_name -> shap: 0.0003  
99 - longitude lga\_airport\_distance\_km -> shap: 0.0003  
100 - total\_min\_cost rate\_adj\_in\_name -> shap: 0.0003  
101 - latitude total\_min\_cost -> shap: 0.0003  
102 - longitude rate\_adj\_in\_name -> shap: 0.0003  
103 - availability\_365 -> shap: 0.0003  
104 - calculated\_host\_listings\_count rate\_adj\_in\_name -> shap: 0.0003  
105 - rate\_nouns\_in\_name rate\_propn\_in\_name -> shap: 0.0003  
106 - price jfk\_airport\_distance\_km -> shap: 0.0003  
107 - jfk\_airport\_distance\_km nb\_propn\_in\_name -> shap: 0.0003  
108 - longitude nb\_propn\_in\_name -> shap: 0.0003  
109 - len\_name nb\_adj\_in\_name -> shap: 0.0003  
110 - total\_min\_cost nb\_nouns\_in\_name -> shap: 0.0003  
111 - latitude price -> shap: 0.0003  
112 - len\_name nb\_nouns\_in\_name -> shap: 0.0003

113 - price lga\_airport\_distance\_km -> shap: 0.0003  
114 - longitude nb\_adv\_in\_name -> shap: 0.0003  
115 - price rate\_propn\_in\_name -> shap: 0.0002  
116 - availability\_365^2 -> shap: 0.0002  
117 - nyc\_tm\_distance\_km avg\_airport\_distance\_km -> shap: 0.0002  
118 - nyc\_cp\_distance\_km rate\_adj\_in\_name -> shap: 0.0002  
119 - longitude ewr\_airport\_distance\_km -> shap: 0.0002  
120 - nyc\_cp\_distance\_km ewr\_airport\_distance\_km -> shap: 0.0002  
121 - nb\_propn\_in\_name rate\_adv\_in\_name -> shap: 0.0002  
122 - price len\_name -> shap: 0.0002  
123 - price nyc\_fd\_distance\_km -> shap: 0.0002  
124 - longitude nyc\_fd\_distance\_km -> shap: 0.0002  
125 - total\_min\_cost rate\_nouns\_in\_name -> shap: 0.0002  
126 - nyc\_fd\_distance\_km avg\_airport\_distance\_km -> shap: 0.0002  
127 - price nb\_propn\_in\_name -> shap: 0.0002  
128 - nb\_propn\_in\_name rate\_nouns\_in\_name -> shap: 0.0002  
129 - nyc\_cp\_distance\_km lga\_airport\_distance\_km -> shap: 0.0002  
130 - calculated\_host\_listings\_count rate\_adv\_in\_name -> shap: 0.0002  
131 - latitude nb\_nouns\_in\_name -> shap: 0.0002  
132 - minimum\_nights nyc\_cp\_distance\_km -> shap: 0.0002  
133 - latitude nb\_propn\_in\_name -> shap: 0.0002  
134 - minimum\_nights rate\_propn\_in\_name -> shap: 0.0002  
135 - total\_min\_cost rate\_adv\_in\_name -> shap: 0.0002  
136 - avg\_airport\_distance\_km nb\_adj\_in\_name -> shap: 0.0002  
137 - calculated\_host\_listings\_count nb\_adj\_in\_name -> shap: 0.0002  
138 - total\_min\_cost^2 -> shap: 0.0002  
139 - longitude avg\_airport\_distance\_km -> shap: 0.0002  
140 - lga\_airport\_distance\_km nb\_adj\_in\_name -> shap: 0.0002  
141 - nb\_adj\_in\_name rate\_nouns\_in\_name -> shap: 0.0002  
142 - ewr\_airport\_distance\_km nb\_nouns\_in\_name -> shap: 0.0002  
143 - ewr\_airport\_distance\_km avg\_airport\_distance\_km -> shap: 0.0002  
144 - nyc\_tm\_distance\_km nb\_nouns\_in\_name -> shap: 0.0002  
145 - len\_name nb\_adv\_in\_name -> shap: 0.0002  
146 - ewr\_airport\_distance\_km nb\_adv\_in\_name -> shap: 0.0001  
147 - nyc\_tm\_distance\_km -> shap: 0.0001  
148 - nb\_adv\_in\_name rate\_nouns\_in\_name -> shap: 0.0001  
149 - total\_min\_cost rate\_propn\_in\_name -> shap: 0.0001  
150 - minimum\_nights rate\_nouns\_in\_name -> shap: 0.0001  
151 - lga\_airport\_distance\_km rate\_propn\_in\_name -> shap: 0.0001  
152 - longitude nb\_nouns\_in\_name -> shap: 0.0001  
153 - price rate\_nouns\_in\_name -> shap: 0.0001  
154 - price nyc\_cp\_distance\_km -> shap: 0.0001  
155 - jfk\_airport\_distance\_km rate\_propn\_in\_name -> shap: 0.0001  
156 - longitude -> shap: 0.0001  
157 - avg\_airport\_distance\_km rate\_propn\_in\_name -> shap: 0.0001  
158 - lga\_airport\_distance\_km nb\_propn\_in\_name -> shap: 0.0001  
159 - calculated\_host\_listings\_count rate\_nouns\_in\_name -> shap: 0.0001  
160 - nyc\_cp\_distance\_km len\_name -> shap: 0.0001  
161 - latitude ewr\_airport\_distance\_km -> shap: 0.0001  
162 - nb\_propn\_in\_name rate\_propn\_in\_name -> shap: 0.0001  
163 - total\_min\_cost nb\_adj\_in\_name -> shap: 0.0001  
164 - nyc\_cp\_distance\_km nyc\_fd\_distance\_km -> shap: 0.0001  
165 - availability\_365 nyc\_cp\_distance\_km -> shap: 0.0001  
166 - jfk\_airport\_distance\_km -> shap: 0.0001  
167 - latitude -> shap: 0.0001  
168 - nyc\_cp\_distance\_km nb\_nouns\_in\_name -> shap: 0.0001

169 - rate\_adj\_in\_name rate\_nouns\_in\_name -> shap: 0.0001  
170 - longitude -> shap: 0.0001  
171 - jfk\_airport\_distance\_km nb\_nouns\_in\_name -> shap: 0.0001  
172 - price -> shap: 0.0001  
173 - nyc\_fd\_distance\_km rate\_nouns\_in\_name -> shap: 0.0001  
174 - longitude^2 -> shap: 0.0001  
175 - nb\_nouns\_in\_name rate\_propn\_in\_name -> shap: 0.0001  
176 - latitude nyc\_fd\_distance\_km -> shap: 0.0001  
177 - lga\_airport\_distance\_km nb\_nouns\_in\_name -> shap: 0.0001  
178 - jfk\_airport\_distance\_km rate\_adv\_in\_name -> shap: 0.0001  
179 - lga\_airport\_distance\_km -> shap: 0.0001  
180 - nyc\_cp\_distance\_km avg\_airport\_distance\_km -> shap: 0.0001  
181 - longitude rate\_propn\_in\_name -> shap: 0.0001  
182 - latitude nb\_adj\_in\_name -> shap: 0.0001  
183 - latitude jfk\_airport\_distance\_km -> shap: 0.0001  
184 - latitude lga\_airport\_distance\_km -> shap: 0.0001  
185 - total\_min\_cost -> shap: 0.0001  
186 - latitude nyc\_cp\_distance\_km -> shap: 0.0001  
187 - lga\_airport\_distance\_km rate\_nouns\_in\_name -> shap: 0.0001  
188 - jfk\_airport\_distance\_km nb\_adj\_in\_name -> shap: 0.0001  
189 - total\_min\_cost nb\_adv\_in\_name -> shap: 0.0001  
190 - total\_min\_cost nb\_propn\_in\_name -> shap: 0.0001  
191 - lga\_airport\_distance\_km ewr\_airport\_distance\_km -> shap: 0.0001  
192 - avg\_airport\_distance\_km rate\_nouns\_in\_name -> shap: 0.0001  
193 - ewr\_airport\_distance\_km rate\_adv\_in\_name -> shap: 0.0001  
194 - price nb\_nouns\_in\_name -> shap: 0.0001  
195 - lga\_airport\_distance\_km nb\_adv\_in\_name -> shap: 0.0001  
196 - price nb\_adv\_in\_name -> shap: 0.0001  
197 - nyc\_tm\_distance\_km nb\_adj\_in\_name -> shap: 0.0001  
198 - longitude price -> shap: 0.0001  
199 - nb\_nouns\_in\_name rate\_adj\_in\_name -> shap: 0.0001  
200 - nyc\_cp\_distance\_km rate\_nouns\_in\_name -> shap: 0.0001  
201 - rate\_adj\_in\_name rate\_propn\_in\_name -> shap: 0.0001  
202 - ewr\_airport\_distance\_km^2 -> shap: 0.0001  
203 - nyc\_fd\_distance\_km rate\_adv\_in\_name -> shap: 0.0001  
204 - nyc\_tm\_distance\_km^2 -> shap: 0.0001  
205 - jfk\_airport\_distance\_km rate\_adj\_in\_name -> shap: 0.0001  
206 - latitude nyc\_tm\_distance\_km -> shap: 0.0001  
207 - nb\_adj\_in\_name nb\_nouns\_in\_name -> shap: 0.0001  
208 - minimum\_nights nb\_propn\_in\_name -> shap: 0.0001  
209 - minimum\_nights nb\_nouns\_in\_name -> shap: 0.0001  
210 - avg\_airport\_distance\_km rate\_adj\_in\_name -> shap: 0.0001  
211 - nb\_propn\_in\_name rate\_adj\_in\_name -> shap: 0.0001  
212 - longitude rate\_nouns\_in\_name -> shap: 0.0001  
213 - jfk\_airport\_distance\_km -> shap: 0.0001  
214 - nb\_adj\_in\_name nb\_propn\_in\_name -> shap: 0.0001  
215 - nb\_adv\_in\_name rate\_propn\_in\_name -> shap: 0.0001  
216 - price rate\_adj\_in\_name -> shap: 0.0001  
217 - price rate\_adv\_in\_name -> shap: 0.0001  
218 - nyc\_fd\_distance\_km nb\_adv\_in\_name -> shap: 0.0001  
219 - price nb\_adj\_in\_name -> shap: 0.0001  
220 - nb\_adv\_in\_name nb\_propn\_in\_name -> shap: 0.0  
221 - nyc\_fd\_distance\_km nb\_adj\_in\_name -> shap: 0.0  
222 - nyc\_fd\_distance\_km nb\_nouns\_in\_name -> shap: 0.0  
223 - calculated\_host\_listings\_count nb\_nouns\_in\_name -> shap: 0.0  
224 - latitude^2 -> shap: 0.0

225 - longitude nb\_adj\_in\_name -> shap: 0.0  
226 - avg\_airport\_distance\_km^2 -> shap: 0.0  
227 - latitude rate\_propn\_in\_name -> shap: 0.0  
228 - ewr\_airport\_distance\_km -> shap: 0.0  
229 - nyc\_fd\_distance\_km rate\_adj\_in\_name -> shap: 0.0  
230 - rate\_adv\_in\_name rate\_nouns\_in\_name -> shap: 0.0  
231 - lga\_airport\_distance\_km rate\_adv\_in\_name -> shap: 0.0  
232 - nyc\_tm\_distance\_km -> shap: 0.0  
233 - nyc\_fd\_distance\_km ewr\_airport\_distance\_km -> shap: 0.0  
234 - nb\_adj\_in\_name nb\_adv\_in\_name -> shap: 0.0  
235 - longitude rate\_adv\_in\_name -> shap: 0.0  
236 - lga\_airport\_distance\_km^2 -> shap: 0.0  
237 - nb\_adv\_in\_name rate\_adj\_in\_name -> shap: 0.0  
238 - minimum\_nights rate\_adv\_in\_name -> shap: 0.0  
239 - jfk\_airport\_distance\_km rate\_nouns\_in\_name -> shap: 0.0  
240 - jfk\_airport\_distance\_km nb\_adv\_in\_name -> shap: 0.0  
241 - nyc\_fd\_distance\_km -> shap: 0.0  
242 - nyc\_cp\_distance\_km -> shap: 0.0  
243 - ewr\_airport\_distance\_km rate\_nouns\_in\_name -> shap: 0.0  
244 - ewr\_airport\_distance\_km -> shap: 0.0  
245 - len\_name -> shap: 0.0  
246 - price^2 -> shap: 0.0  
247 - lga\_airport\_distance\_km -> shap: 0.0  
248 - len\_name^2 -> shap: 0.0  
249 - nb\_adj\_in\_name rate\_adj\_in\_name -> shap: 0.0  
250 - rate\_adj\_in\_name rate\_adv\_in\_name -> shap: 0.0  
251 - nb\_nouns\_in\_name rate\_adv\_in\_name -> shap: 0.0  
252 - rate\_propn\_in\_name -> shap: 0.0  
253 - nyc\_cp\_distance\_km rate\_adv\_in\_name -> shap: 0.0  
254 - avg\_airport\_distance\_km nb\_nouns\_in\_name -> shap: 0.0  
255 - lga\_airport\_distance\_km rate\_adj\_in\_name -> shap: 0.0  
256 - avg\_airport\_distance\_km nb\_propn\_in\_name -> shap: 0.0  
257 - nyc\_tm\_distance\_km rate\_adv\_in\_name -> shap: 0.0  
258 - latitude rate\_nouns\_in\_name -> shap: 0.0  
259 - latitude -> shap: 0.0  
260 - nb\_nouns\_in\_name rate\_nouns\_in\_name -> shap: 0.0  
261 - longitude nyc\_cp\_distance\_km -> shap: 0.0  
262 - rate\_adj\_in\_name^2 -> shap: 0.0  
263 - availability\_365 nb\_adj\_in\_name -> shap: 0.0  
264 - price minimum\_nights -> shap: 0.0  
265 - nb\_nouns\_in\_name -> shap: 0.0  
266 - ewr\_airport\_distance\_km nb\_adj\_in\_name -> shap: 0.0  
267 - avg\_airport\_distance\_km -> shap: 0.0  
268 - nyc\_fd\_distance\_km^2 -> shap: 0.0  
269 - nyc\_cp\_distance\_km -> shap: 0.0  
270 - rate\_adj\_in\_name -> shap: 0.0  
271 - rate\_nouns\_in\_name -> shap: 0.0  
272 - nyc\_tm\_distance\_km rate\_nouns\_in\_name -> shap: 0.0  
273 - rate\_propn\_in\_name -> shap: 0.0  
274 - jfk\_airport\_distance\_km^2 -> shap: 0.0  
275 - rate\_nouns\_in\_name -> shap: 0.0  
276 - avg\_airport\_distance\_km -> shap: 0.0  
277 - rate\_propn\_in\_name^2 -> shap: 0.0  
278 - minimum\_nights -> shap: 0.0  
279 - rate\_adv\_in\_name^2 -> shap: 0.0  
280 - nb\_adv\_in\_name nb\_nouns\_in\_name -> shap: 0.0

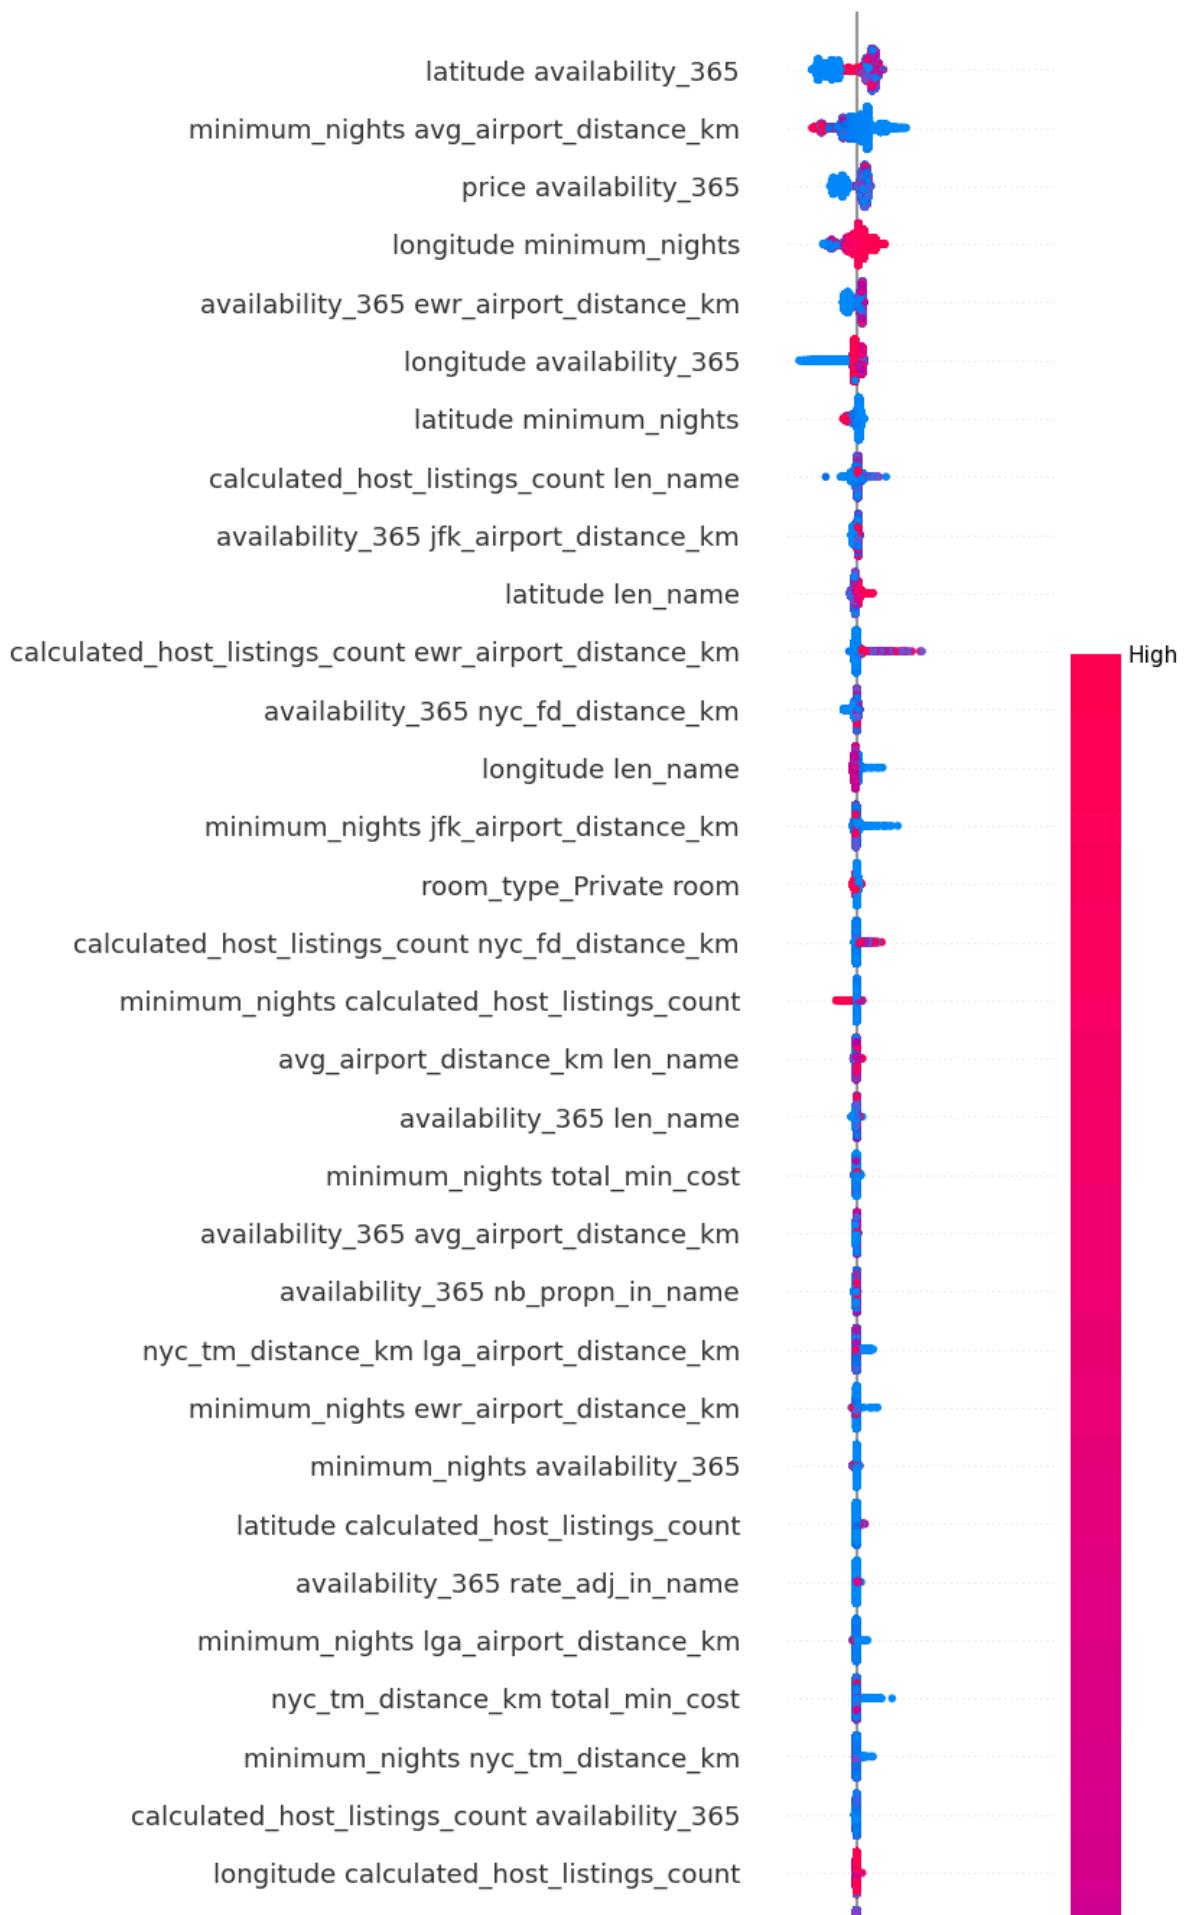
```
281 - len_name rate_nouns_in_name -> shap: 0.0
282 - latitude rate_adv_in_name -> shap: 0.0
283 - nyc_cp_distance_km nb_adv_in_name -> shap: 0.0
284 - calculated_host_listings_count^2 -> shap: 0.0
285 - nb_propn_in_name -> shap: 0.0
286 - nb_adj_in_name rate_propn_in_name -> shap: 0.0
287 - rate_adj_in_name -> shap: 0.0
288 - avg_airport_distance_km rate_adv_in_name -> shap: 0.0
289 - nb_propn_in_name -> shap: 0.0
290 - neighbourhood_group_Manhattan -> shap: 0.0
291 - nyc_fd_distance_km -> shap: 0.0
292 - rate_nouns_in_name^2 -> shap: 0.0
293 - nyc_cp_distance_km^2 -> shap: 0.0
294 - nb_adj_in_name -> shap: 0.0
295 - total_min_cost -> shap: 0.0
296 - len_name rate_propn_in_name -> shap: 0.0
297 - nb_adj_in_name^2 -> shap: 0.0
298 - ewr_airport_distance_km rate_adj_in_name -> shap: 0.0
299 - avg_airport_distance_km nb_adv_in_name -> shap: 0.0
300 - len_name rate_adj_in_name -> shap: 0.0
301 - minimum_nights nb_adv_in_name -> shap: 0.0
302 - calculated_host_listings_count -> shap: 0.0
303 - len_name rate_adv_in_name -> shap: 0.0
304 - rate_adv_in_name -> shap: 0.0
305 - nb_adv_in_name -> shap: 0.0
306 - nyc_tm_distance_km nb_adv_in_name -> shap: 0.0
307 - calculated_host_listings_count -> shap: 0.0
308 - len_name -> shap: 0.0
309 - price -> shap: 0.0
310 - latitude nb_adv_in_name -> shap: 0.0
311 - minimum_nights nb_adj_in_name -> shap: 0.0
312 - 1 -> shap: 0.0
313 - minimum_nights -> shap: 0.0
314 - nb_adv_in_name -> shap: 0.0
315 - nb_nouns_in_name -> shap: 0.0
316 - minimum_nights^2 -> shap: 0.0
317 - calculated_host_listings_count nb_adv_in_name -> shap: 0.0
318 - availability_365 nb_adv_in_name -> shap: 0.0
319 - availability_365 nb_nouns_in_name -> shap: 0.0
320 - availability_365 rate_adv_in_name -> shap: 0.0
321 - nb_adj_in_name rate_adv_in_name -> shap: 0.0
322 - nb_adv_in_name^2 -> shap: 0.0
323 - nb_adv_in_name rate_adv_in_name -> shap: 0.0
324 - nb_nouns_in_name^2 -> shap: 0.0
325 - nb_propn_in_name^2 -> shap: 0.0
326 - nb_adj_in_name -> shap: 0.0
327 - rate_adv_in_name -> shap: 0.0
328 - neighbourhood_group_Brooklyn -> shap: 0.0
329 - neighbourhood_group_Queens -> shap: 0.0
330 - neighbourhood_group_Staten Island -> shap: 0.0
331 - room_type_Shared room -> shap: 0.0
```

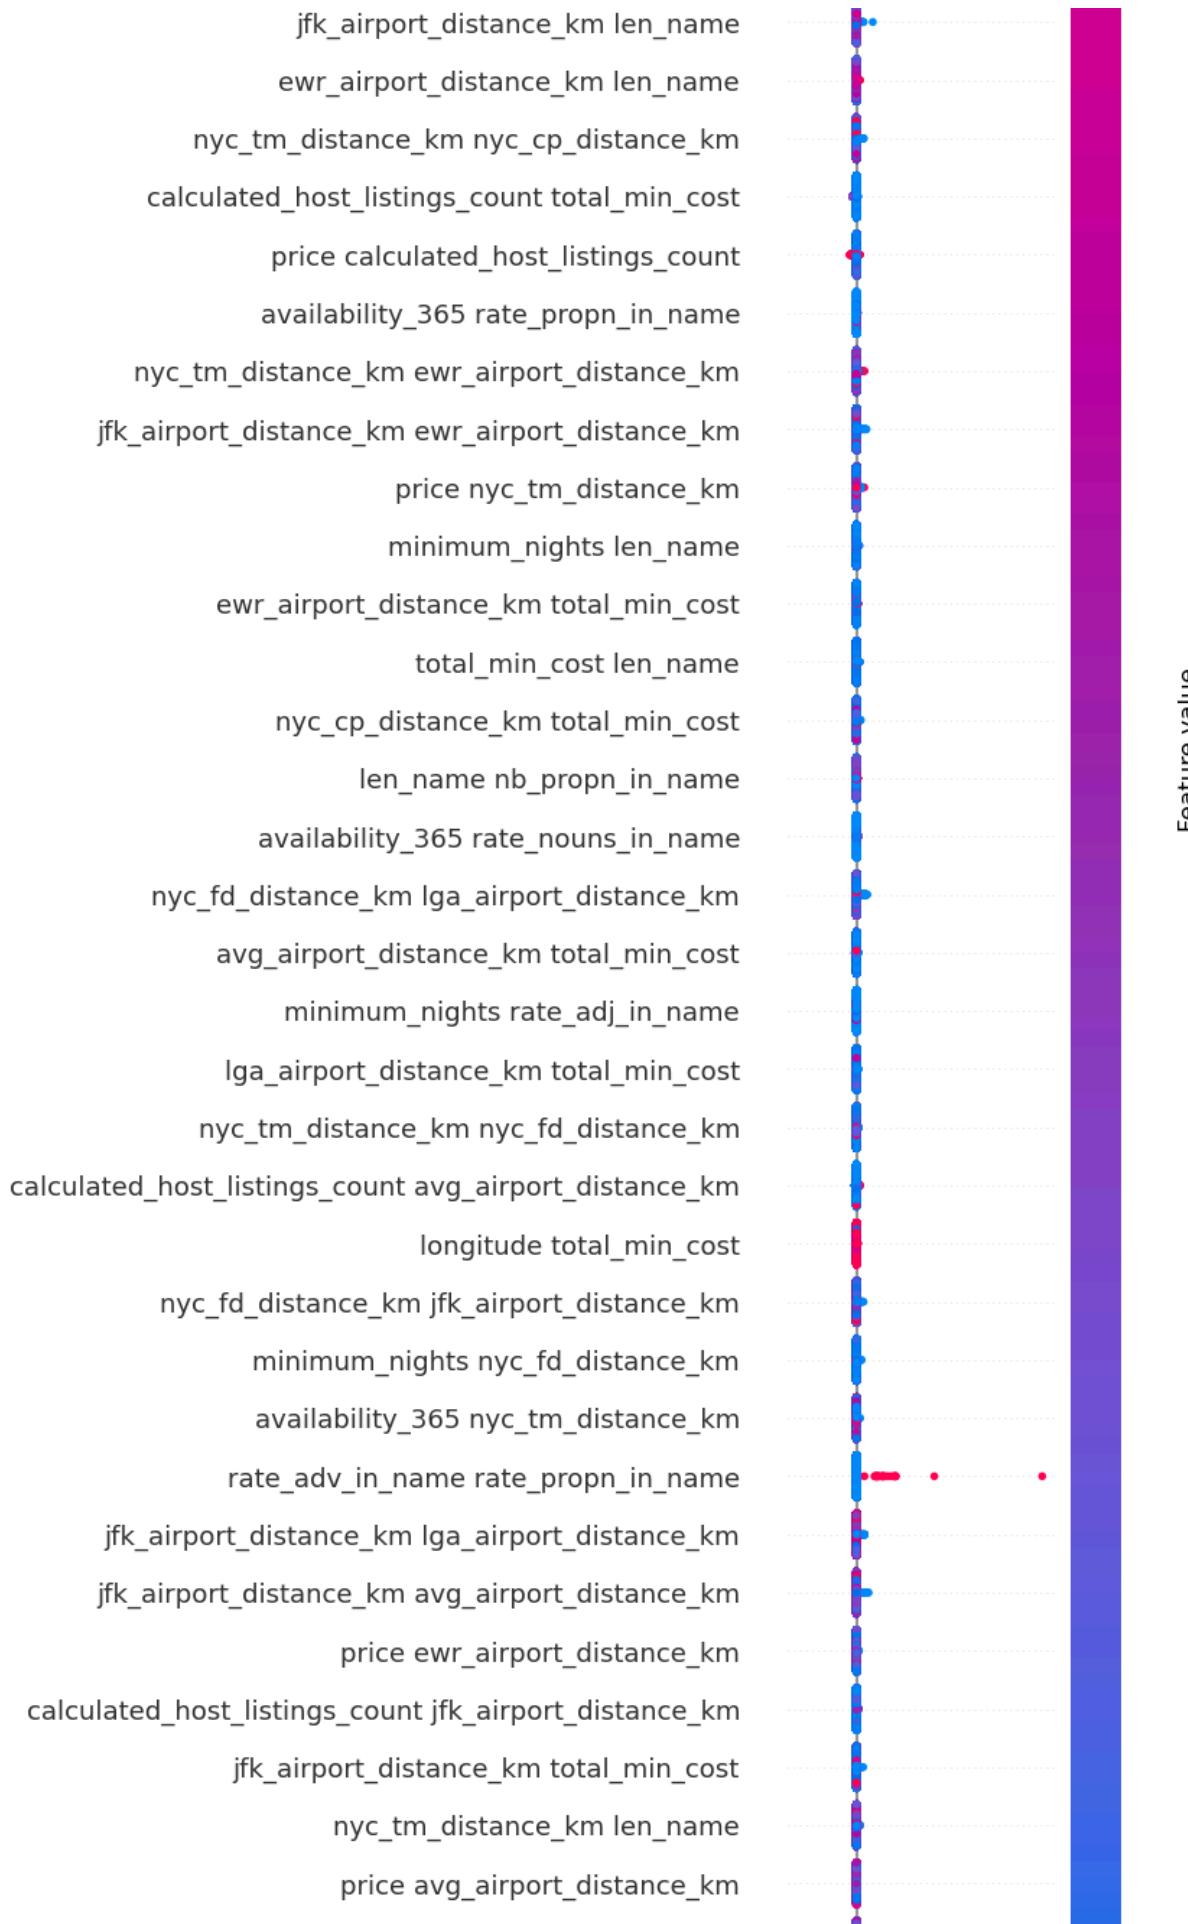
```
In [ ]: plt.rcParams.update({
    "font.size": 8,
    "axes.titlesize": 9,
    "axes.labelsize": 8,
```

```
    "xtick.labelsize": 7,  
    "ytick.labelsize": 7,  
})  
  
shap.summary_plot(  
    shap_values,  
    X_train_soft_s,  
    feature_names=clean_features,  
    max_display=len(selected_features),  
)
```

/var/folders/gc/x3q0m4h564d61kh41xdz\_gy00000gn/T/ipykernel\_4013/4231117112.py:1: FutureWarning: The NumPy global RNG was seeded by calling `np.random.seed`. In a future version this function will no longer use the global RNG. Pass `rng` explicitly to opt-in to the new behaviour and silence this warning.

```
    shap.summary_plot(shap_values, X_train_soft_s, feature_names=clean_features, max_display=len(selected_features))
```







We used SHAP for feature selection as this method is used to measure each feature contribution to predictions (based on a game theory approach model). SHAP accounts for interactions among features (polynomial) and helps to separate valuable features in the model from random noise.

Here we added two random features as a baseline and we expect to have some valuable interactions among features in the model. If there are values lower than noise, is nothing

that can be predicted. If we see the SHAP value graph we can overlook that geographic features as longitude and latitude appear in the most frequent interactions and availability and booking constraints matter (minimum\_nights and availability\_365).

The most important interaction among features are:

1. Latitude \* availability\_365 (0.2634): Most important interaction, so the location combined with how often the listing is available strongly impacts the target.
2. Minimum\_nights \* avg\_airport\_distance\_km (0.2056): Booking requirements and airport accessibility together matter significantly.
3. Price \* availability\_365 (0.1541): Pricing strategy and availability interact to influence outcomes.
4. Longitude \* minimum\_nights (0.1267): East-west location combined with stay requirements.

## 10. Hyperparameter optimization

`rubric={accuracy,reasoning}`

### Your tasks:

Make some attempts to optimize hyperparameters for the models you've tried and summarize your results. In at least one case you should be optimizing multiple hyperparameters for a single model. You may use `sklearn`'s methods for hyperparameter optimization or fancier Bayesian optimization methods. Briefly summarize your results.

- `GridSearchCV`
- `RandomizedSearchCV`
- `scikit-optimize`

Points: 6

Type your answer here, replacing this text.

```
In [68]: param_grid = {
    "histgradientboostingregressor__max_depth": randint(40, 60),
    "histgradientboostingregressor__max_iter": randint(700, 1000),
    "histgradientboostingregressor__learning_rate": loguniform(0.005, 0.015)
}

random_search_hgbr = RandomizedSearchCV(
    make_pipeline(
        general_preprocessor,
        HistGradientBoostingRegressor(
            random_state=RANDOM_STATE
```

```

        )
),
param_distributions=param_grid,
n_iter=5,
n_jobs= -1,
return_train_score=True,
verbose=3
)

random_search_hgbr.fit(X_train, y_train)

print(
    "Random Search best model score: \t %.3f" % random_search_hgbr.best_score_
)

```

pd.DataFrame(random\_search\_hgbr.cv\_results\_)[

- [
- "mean\_test\_score",
- "mean\_train\_score",
- "param\_histgradientboostingregressor\_\_max\_depth",
- "param\_histgradientboostingregressor\_\_max\_iter",
- "param\_histgradientboostingregressor\_\_learning\_rate",
- "mean\_fit\_time",
- "rank\_test\_score",

].set\_index("rank\_test\_score").sort\_index()

Fitting 5 folds for each of 5 candidates, totalling 25 fits  
 Random Search best model score:               0.375

Out[68]:

	mean_test_score	mean_train_score	param_histgradientboostingregressor__m:
rank_test_score			
1	0.374527	0.526722	
2	0.374371	0.535432	
3	0.373817	0.538871	
4	0.373741	0.511925	
5	0.373281	0.541082	

Our hyperparameter optimization is improving model performance. The best model (GBR) improves from R squared of 0.369 to 0.375 after tuning , a modest but meaningful gain.

The train score of 0.527 indicates a reasonable overfitting gap of 0.15. The optimal configuration uses moderately deep trees (max\_depth = 24) with a low learning rate (0.006), balancing model complexity with generalization. So, the model seem stable with a good hyperparameter selection.

In [ ]: ss.

```
In [ ]: ...
```

## 11. Interpretation and feature importances

`rubric={accuracy,reasoning}`

**Your tasks:**

1. Use the methods we saw in class (e.g., `permutation_importance` or `shap`) (or any other methods of your choice) to examine the most important features of one of the non-linear models.
2. Summarize your observations.

*Points:* 8

Using SHAP values as explained before we can summarize the features that have the most significant impact:

- Price \*\* availability\_365: Most important feature as the interaction between listing price and how often it's available has the strongest impact on predictions. Both high and low values influence the model.
- Longitude \*\* minimum\_nights: It seems that geographic location combined with booking requirements is highly influential.
- Availability\_365 \*\* airport distances. It seems that how often a listing is available, combined with its distance from airports/landmarks, matters significantly.
- Latitude \*\* minimum\_nights: North-south location interacting with stay requirements.

We can also see that availability is all over the place in terms of appearing in many features, so this feature seems crucial. It is also noted that geography is not enough to have a good predictive power, we need a set of combinations or interactions with other valuable features as price or availability in order to have better power or prediction.

And also and very important, we can see that the polynomial transformations have better coefficients than single features, so we can almost be certain (as Ridge model also it does not perform well) that there is not a linear relationship here.

```
In [ ]: # For random forest!
```

```
# explainer = shap.TreeExplainer(tr)
# shap_values = explainer.shap_values(X_train_soft_s)
# shap.summary_plot(
#     shap_values,
#     X_train_soft_s,
#     feature_names=clean_features,
#     max_display=len(selected_features),
# )
```

```
In [75]: from scipy import sparse

best_pipe = random_search_hgbr.best_estimator_
preprocess = best_pipe[:-1]
model = best_pipe.named_steps["histgradientboostingregressor"]

n = min(2000, len(X_train))
X_shap = X_train.sample(n=n, random_state=RANDOM_STATE)

X_shap_t = preprocess.transform(X_shap)
feature_names = preprocess.get_feature_names_out()

if sparse.issparse(X_shap_t):
    X_shap_t = X_shap_t.toarray()

X_shap_df = pd.DataFrame(X_shap_t, columns=feature_names)

explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(X_shap_df)

plt.rcParams.update({
    "font.size": 3,
    "axes.titlesize": 3,
    "axes.labelsize": 3,
    "xtick.labelsize": 3,
    "ytick.labelsize": 3,
})

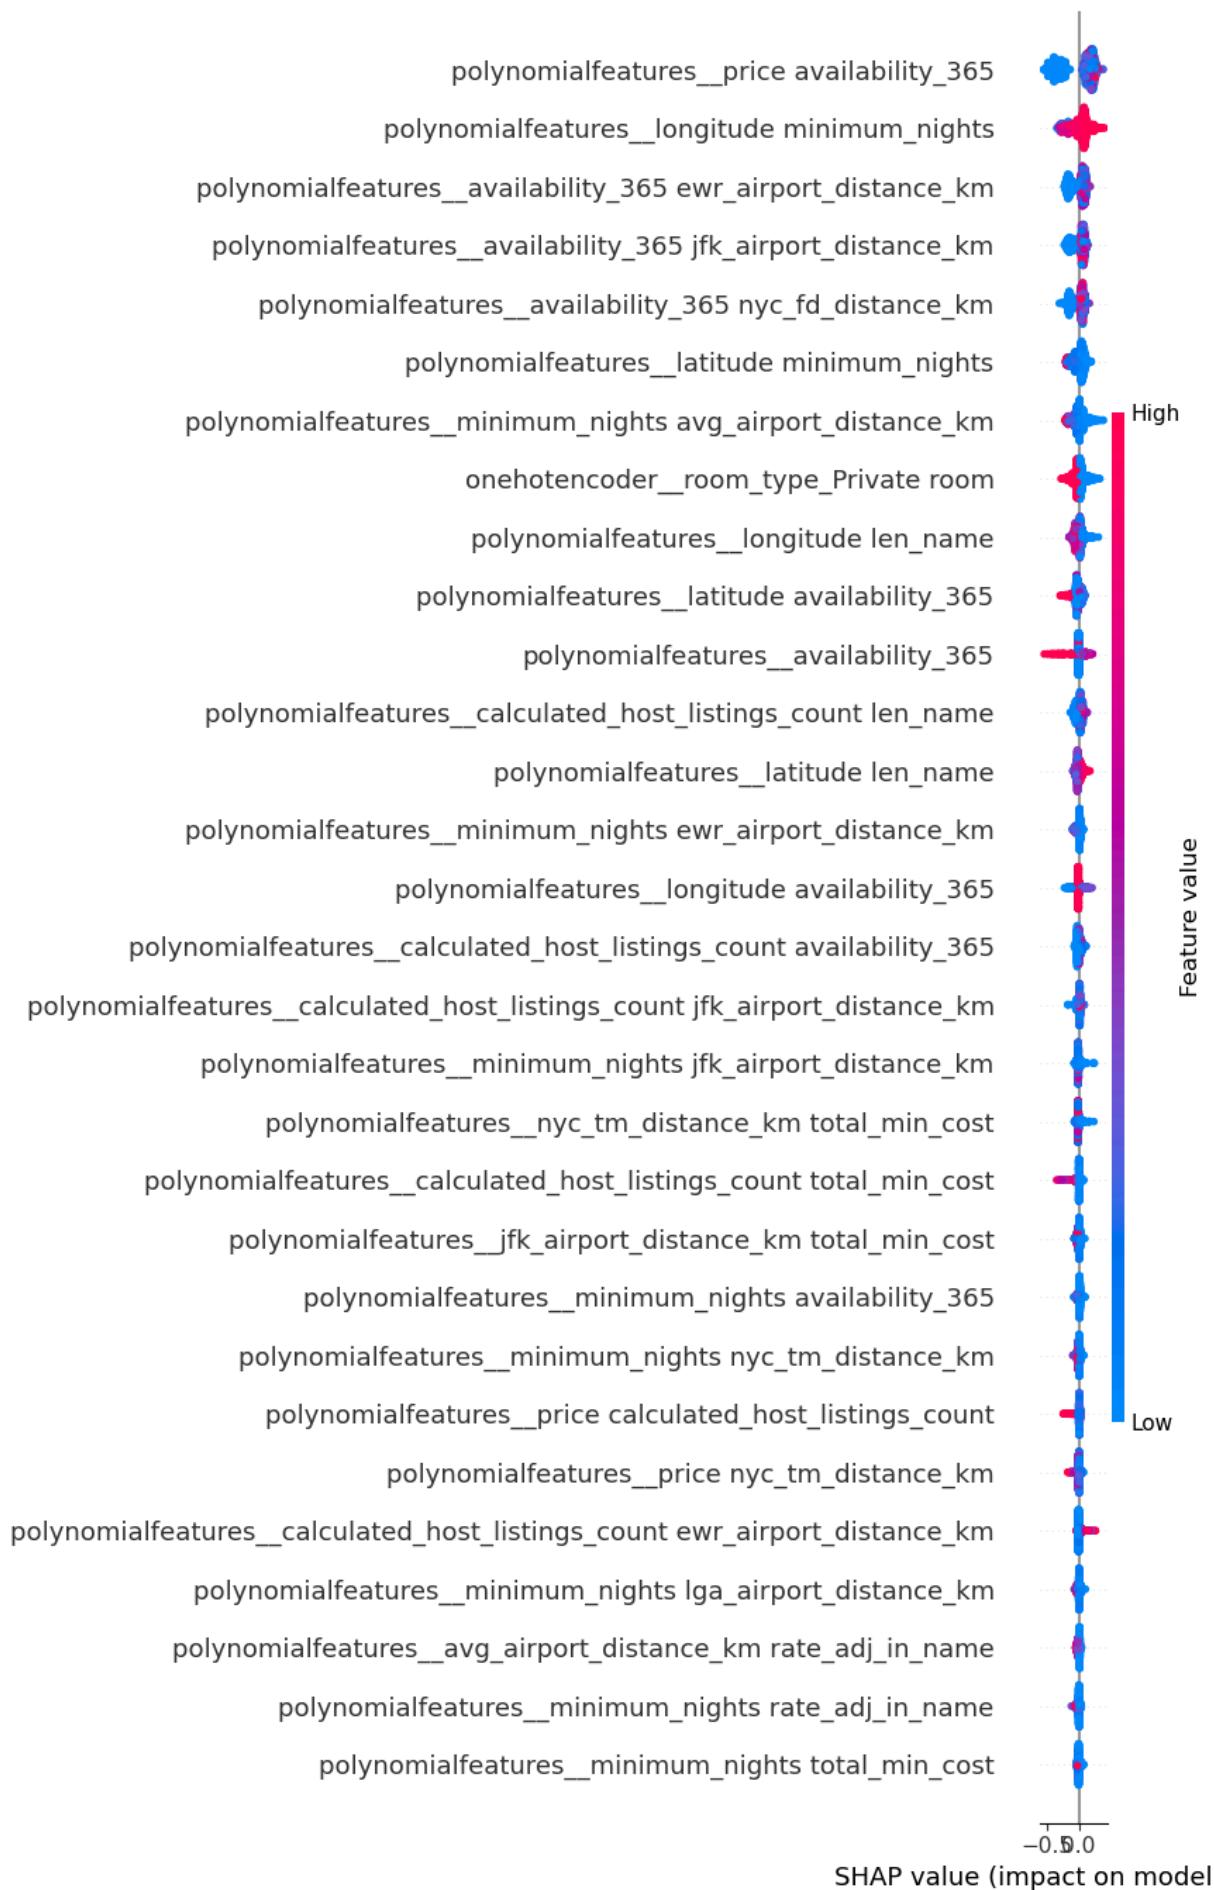
shap.summary_plot(
    shap_values,
    X_shap_df,
    max_display=30
)
plt.tight_layout()
plt.show()

shap.summary_plot(
    shap_values,
    X_shap_df,
    plot_type="bar",
```

```
    max_display=30
)
plt.tight_layout()
plt.show()
```

C:\Users\hpala\AppData\Local\Temp\ipykernel\_25104\2574226130.py:35: FutureWarning: The NumPy global RNG was seeded by calling `np.random.seed`. In a future version this function will no longer use the global RNG. Pass `rng` explicitly to opt-in to the new behaviour and silence this warning.

```
shap.summary_plot(
```

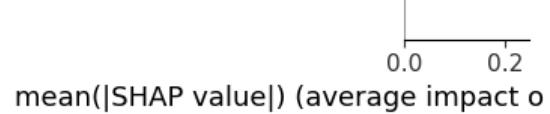


```
<Figure size 640x480 with 0 Axes>
```

```
C:\Users\hpala\AppData\Local\Temp\ipykernel_25104\2574226130.py:44: FutureWarning: T  
he NumPy global RNG was seeded by calling `np.random.seed`. In a future version this  
function will no longer use the global RNG. Pass `rng` explicitly to opt-in to the n  
ew behaviour and silence this warning.
```

```
shap.summary_plot(
```





<Figure size 640x480 with 0 Axes>

In [ ]:

In [ ]: ...

## 12. Results on the test set

rubric={accuracy,reasoning}

### Your tasks:

1. Try your best performing model on the test data and report test scores.
2. Do the test scores agree with the validation scores from before? To what extent do you trust your results? Do you think you've had issues with optimization bias?
3. Take one or two test predictions and explain them with SHAP force plots.

Points: 6

We have a test score of 0.3648

In a comparison with validation:

- Validation: 0.3745 (from hyperparameter tuning)
- Test: 0.3648
- Difference: 0.0097

The test score closely matches the validation score, with only a 2.6% relative decrease. This strong agreement indicates:

1. **Minimal optimization bias:** Our hyperparameter search didn't overfit to the validation set
2. **Good generalization:** The model performs consistently on unseen data
3. **Trustworthy results:** The 37% variance explained represents genuine predictive capability

The modest absolute R squared (0.37) reflects the inherent difficulty of predicting subjective review scores from listing features alone. Unmeasured factors (host responsiveness, cleanliness, guest expectations) contribute to the remaining 63% of variance. Given these limitations, capturing 37% of variance is reasonable performance for this prediction task.

The prediction models tells us that expensive and remote listings receive poor reviews. Despite its high total\_min\_cost (160,000) and price × total\_min\_cost (32,000), the model predicts a review score of only 0.30. The combination of high costs, inconvenient airport access (high distance interactions), and remote location creates guest dissatisfaction. This demonstrates the model learned that expensive listings in less accessible locations tend to disappoint guests. And we have confirmation that price alone doesn't guarantee good reviews; location and accessibility matter significantly.

SHAP VALUES: This listing has much lower predicted value (0.3 vs baseline 1.336) primarily because its price-availability interaction is zero, and several airport distance interactions are also zero or negative. This suggests it might be a listing with low availability or unusual pricing.

```
In [74]: best_pipe.score(X_test, y_test)
```

```
Out[74]: 0.3648321574290111
```

```
In [76]: ex_index = 573

best_pipe = random_search_hgbr.best_estimator_
preprocess = best_pipe[:-1]
model = best_pipe.named_steps["histgradientboostingregressor"]

x_raw = X_test.iloc[[ex_index]]
x_enc = preprocess.transform(x_raw)

feature_names = preprocess.get_feature_names_out()

if sparse.issparse(x_enc):
    x_enc = x_enc.toarray()

x_enc_df = pd.DataFrame(x_enc, columns=feature_names)

display(x_enc_df.iloc[0].sort_values(ascending=False).head(30))

pred = best_pipe.predict(x_raw)[0]
print("Prediction:", pred)

explainer = shap.TreeExplainer(model)

shap_values = explainer.shap_values(x_enc_df)

shap_df = (pd.DataFrame(
    {"feature": feature_names, "shap_value": shap_values[0], "value": x_enc_df.iloc[0].abs_shap=lambda d: d.shap_value.abs()})
```

```

    .sort_values("abs_shap", ascending=False))

display(shap_df[["feature", "value", "shap_value"]].head(30))

shap.initjs()

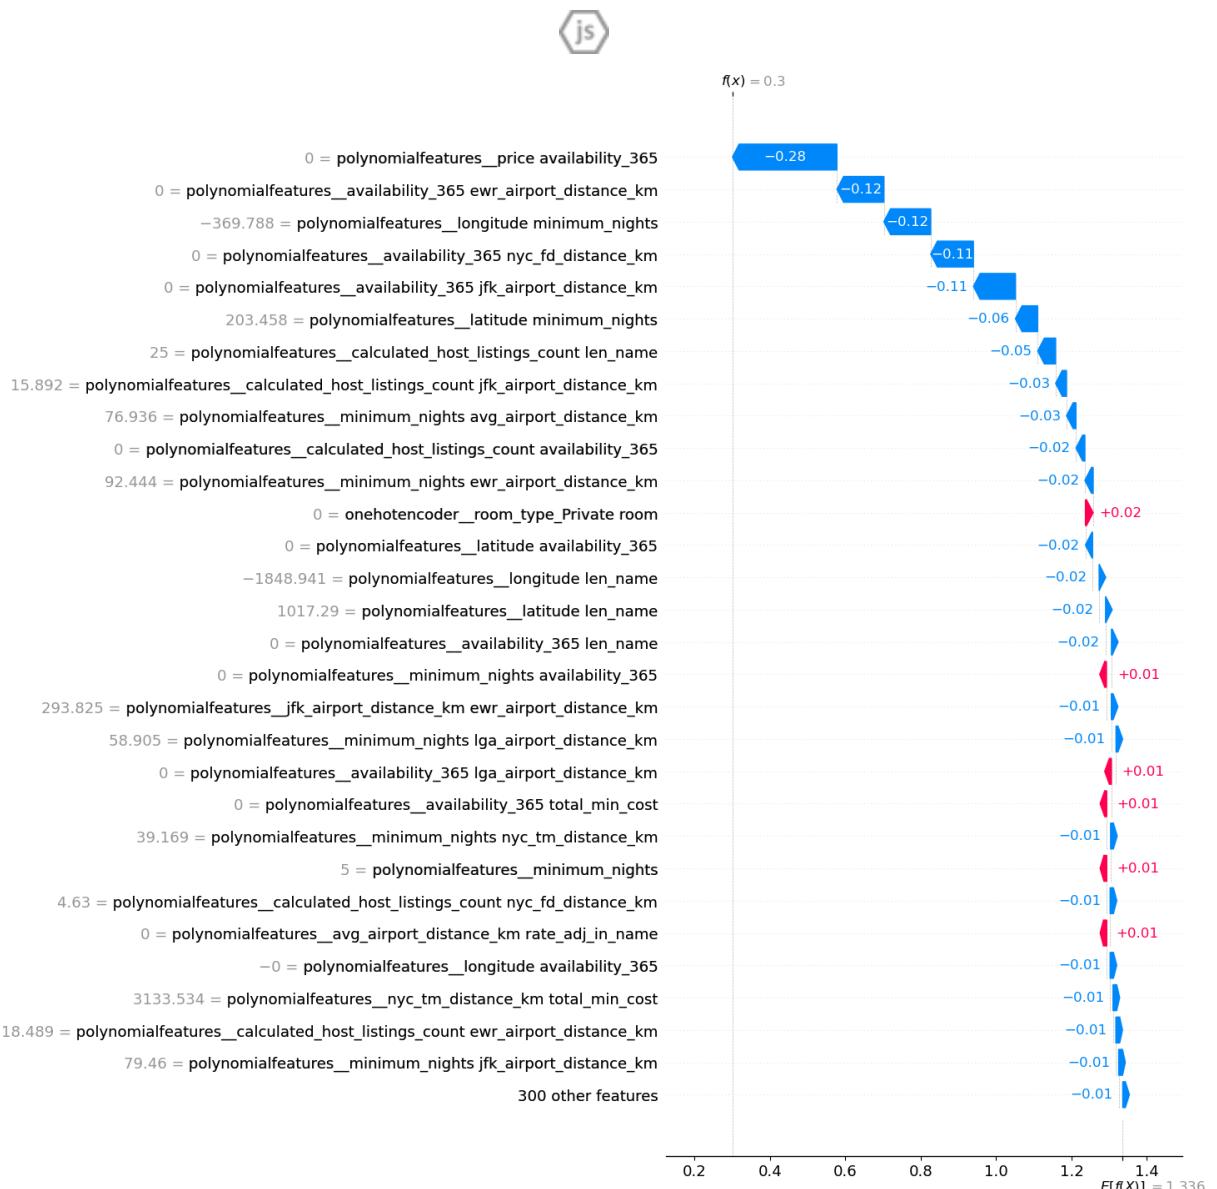
base_value = explainer.expected_value
exp = shap.Explanation(
    values=shap_values[0],
    base_values=base_value,
    data=x_enc_df.iloc[0].values,
    feature_names=feature_names,
)

shap.plots.waterfall(exp, max_display=30)

```

polynomialfeatures__total_min_cost^2	160000.000000
polynomialfeatures__price total_min_cost	32000.000000
polynomialfeatures__latitude total_min_cost	16276.636000
polynomialfeatures__total_min_cost len_name	10000.000000
polynomialfeatures__ewr_airport_distance_km total_min_cost	7395.550686
polynomialfeatures__price^2	6400.000000
polynomialfeatures__jfk_airport_distance_km total_min_cost	6356.802021
polynomialfeatures__avg_airport_distance_km total_min_cost	6154.918861
polynomialfeatures__longitude^2	5469.733994
polynomialfeatures__lga_airport_distance_km total_min_cost	4712.403876
polynomialfeatures__nyc_cp_distance_km total_min_cost	4056.797245
polynomialfeatures__latitude price	3255.327200
polynomialfeatures__nyc_tm_distance_km total_min_cost	3133.534280
polynomialfeatures__price len_name	2000.000000
polynomialfeatures__minimum_nights total_min_cost	2000.000000
polynomialfeatures__nyc_fd_distance_km total_min_cost	1852.158863
polynomialfeatures__latitude^2	1655.805497
polynomialfeatures__price ewr_airport_distance_km	1479.110137
polynomialfeatures__price jfk_airport_distance_km	1271.360404
polynomialfeatures__price avg_airport_distance_km	1230.983772
polynomialfeatures__total_min_cost nb_propn_in_name	1200.000000
polynomialfeatures__latitude len_name	1017.289750
polynomialfeatures__price lga_airport_distance_km	942.480775
polynomialfeatures__price nyc_cp_distance_km	811.359449
polynomialfeatures__latitude ewr_airport_distance_km	752.341791
polynomialfeatures__latitude jfk_airport_distance_km	646.670954
polynomialfeatures__price nyc_tm_distance_km	626.706856
polynomialfeatures__latitude avg_airport_distance_km	626.133587
polynomialfeatures__len_name^2	625.000000
polynomialfeatures__latitude lga_airport_distance_km	479.388016
Name: 573, dtype: float64	
Prediction: 0.3003480226733975	

		<b>feature</b>	<b>value</b>	<b>shap_value</b>
72		polynomialfeatures_price availability_365	0.000000	-0.277241
135		polynomialfeatures_availability_365 ewr_airpo...	0.000000	-0.124474
49		polynomialfeatures_longitude minimum_nights	-369.788250	-0.124358
132		polynomialfeatures_availability_365 nyc_fd_di...	0.000000	-0.113238
133		polynomialfeatures_availability_365 jfk_airpo...	0.000000	-0.111357
27		polynomialfeatures_latitude minimum_nights	203.457950	-0.059163
120		polynomialfeatures_calculated_host_listings_c...	25.000000	-0.047551
115		polynomialfeatures_calculated_host_listings_c...	15.892005	-0.028901
99		polynomialfeatures_minimum_nights avg_airport...	76.936486	-0.025167
111		polynomialfeatures_calculated_host_listings_c...	0.000000	-0.023566
98		polynomialfeatures_minimum_nights ewr_airport...	92.444384	-0.021954
327		onehotencoder_room_type_Private room	0.000000	0.020540
29		polynomialfeatures_latitude availability_365	0.000000	-0.019084
60		polynomialfeatures_longitude len_name	-1848.941250	-0.017290
38		polynomialfeatures_latitude len_name	1017.289750	-0.017076
138		polynomialfeatures_availability_365 len_name	0.000000	-0.015625
92		polynomialfeatures_minimum_nights availabilit...	0.000000	0.013395
197		polynomialfeatures_jfk_airport_distance_km ew...	293.825322	-0.013034
97		polynomialfeatures_minimum_nights lga_airport...	58.905048	-0.012519
134		polynomialfeatures_availability_365 lga_airpo...	0.000000	0.012338
137		polynomialfeatures_availability_365 total_min...	0.000000	0.012272
93		polynomialfeatures_minimum_nights nyc_tm_dist...	39.169179	-0.010146
4		polynomialfeatures_minimum_nights	5.000000	0.010106
114		polynomialfeatures_calculated_host_listings_c...	4.630397	-0.008888
241		polynomialfeatures_avg_airport_distance_km ra...	0.000000	0.008823
51		polynomialfeatures_longitude availability_365	-0.000000	-0.008736
154		polynomialfeatures_nyc_tm_distance_km total_m...	3133.534280	-0.007799
117		polynomialfeatures_calculated_host_listings_c...	18.488877	-0.007549
96		polynomialfeatures_minimum_nights jfk_airport...	79.460025	-0.007480
11		polynomialfeatures_lga_airport_distance_km	11.781010	0.007046



In [ ]: ...

In [ ]: ...

In [ ]: ...

In [ ]: ...

## 13. Summary of results

`rubric={reasoning}`

Imagine that you want to present the summary of these results to your boss and co-workers.

**Your tasks:**

1. Create a table summarizing important results.
2. Write concluding remarks.
3. Discuss other ideas that you did not try but could potentially improve the performance/interpretability .
4. Report your final test score along with the metric you used at the top of this notebook.

*Points:* 8

## Model Performance

Our gradient boosting model achieved a test R squared of 0.365, successfully explaining approximately 37% of the variance in Airbnb listing review frequency. While this may appear modest in absolute terms, it represents meaningful predictive capability given the inherent challenges of the problem:

What we have found?

- Strong Generalization: The close alignment between validation (0.375) and test (0.365) scores indicates minimal optimization bias and robust model performance on unseen data.
- Feature Interactions : Polynomial features significantly outperformed single features, demonstrating that review frequency depends on complex interactions rather than linear relationships.

## Model Comparison Insights

1. Ridge regression severely underfit ( $R^2 = 0.253$ ), requiring extreme regularization due to multicollinearity
2. Decision trees catastrophically overfit (perfect train, negative test  $R^2$ )
3. Random Forest showed promise but significant overfitting (0.847 train vs 0.338 test)
4. Gradient Boosting achieved the best balance with modest overfitting (0.527 train vs 0.365 test)

For Airbnb hosts, this model provides some insights:

- Availability is crucial: Listings with higher availability combined with competitive pricing receive more reviews
- Location matters differently than expected: Proximity to airports and landmarks interacts with other features rather than having direct effects
- Booking constraints impact reviews: Minimum night requirements influence review frequency, especially when combined with geographic factors
- Room type has limited direct impact: Being a private room vs entire home shows minimal effect (SHAP: 0.021)

## Model Limitations

The 63% unexplained variance reflects legitimate constraints:

- Missing subjective factors: Host responsiveness, cleanliness, guest experience, accuracy of listing descriptions
- Selection bias in reviews: Dissatisfied guests are more likely to leave reviews than satisfied ones (this we transformed in the beginning of the data in order to prevent data leakage)
- Temporal dynamics: Seasonal patterns, booking trends, and time-varying host behavior not captured
- Guest characteristics: Demographics, travel purpose, and expectations unavailable in the dataset

In [ ]: ...

## 14. Creating a data analysis pipeline (Challenging)

rubric={reasoning}

### Your tasks:

- Convert this notebook into scripts to create a reproducible data analysis pipeline with appropriate documentation. Submit your project folder in addition to this notebook on GitHub and briefly comment on your organization in the text box below.

Points: 0.5

Type your answer here, replacing this text.

## 15. Your takeaway from the course (Challenging)

rubric={reasoning}

### Your tasks:

What is your biggest takeaway from this course?

Points: 0.5

Type your answer here, replacing this text.

### **Restart, run all and export a PDF before submitting**

Before submitting, don't forget to run all cells in your notebook to make sure there are no errors and so that the TAs can see your plots on Gradescope. You can do this by clicking the  button or going to Kernel -> Restart Kernel and Run All Cells... in the menu. This is not only important for MDS, but a good habit you should get into before ever committing a notebook to GitHub, so that your collaborators can run it from top to bottom without issues.

After running all the cells, export a PDF of the notebook (preferably the WebPDF export) and upload this PDF together with the ipynb file to Gradescope (you can select two files when uploading to Gradescope)

---

## Help us improve the labs

The MDS program is continually looking to improve our courses, including lab questions and content. The following optional questions will not affect your grade in any way nor will they be used for anything other than program improvement:

1. Approximately how many hours did you spend working or thinking about this assignment (including lab time)?

## Ans:

2. Do you have any feedback on the lab you be willing to share? For example, any part or question that you particularly liked or disliked?

## Ans:

## Submission

Make sure you have run all cells in your notebook in order before running the cell below, so that all images/graphs appear in the output. The cell below will generate a zip file for you to submit. **Please save before exporting!**

```
In [ ]: # Save your notebook first, then run this cell to export your submission.  
grader.export(run_tests=True)
```